Modelling Diabetes Health Indicators

MA7442 Modelling Data — Assessment 1 (Task 2)

Shahmeer Shahzad

2026-03-02

1. Introduction and Reflection on Task 1 Feedback

This report constitutes Task 2 of the MA7442 *Modelling Data* assessment. It builds directly on the exploratory data analysis (EDA) and research questions developed in Task 1, incorporating the feedback received and extending the analytical scope accordingly.

In Task 1, the EDA was intentionally centred on diabetes status as the primary outcome variable. This focus was adopted to maintain analytical coherence and prevent the report from becoming overly expansive in scope. Given the breadth of variables available in the dataset, conducting a fully comprehensive EDA across multiple outcomes at that stage would likely have diluted the conceptual focus and resulted in excessive length. The analysis was therefore deliberately targeted to support a diabetes-oriented modelling framework.

Feedback indicated, however, that restricting the EDA exclusively to diabetes limited the diversity of subsequent research questions and led to conceptual overlap. Task 2 therefore broadens the analytical framework and introduces dedicated non-diabetes outcome analyses, substantially expanding the empirical foundation of the study.

1.1 Reflection on Task 1 Feedback

Three areas of feedback were identified for improvement.

Scope of the EDA.
The Task 1 EDA effectively characterised the dataset but examined associations solely in relation to diabetes status. This constrained the diversity of the research questions. In Task 2, Section 3 introduces a comprehensive EDA block examining cardiovascular indicators, general health measures, and behavioural variables as primary outcomes in their own right.

Research question overlap.
Because the initial EDA centred exclusively on diabetes, the research questions were conceptually similar. In response, the structure has been revised. RQ1 is retained in refined form; the previous RQ2 and RQ3 have been consolidated into a single modelling-focused question (new RQ2); and two new research questions (RQ3 and RQ4) are derived directly from the expanded non-diabetes EDA.

Figure line type.
In Task 1, a `geom_line` layer was used to connect mean unhealthy days across age groups. Since age group is an ordered categorical variable rather than a continuous measure, this implied unjustified linear interpolation. In Task 2, age-structured visualisations have been revised to use `geom_col` or lollipop-style representations, preserving ordinal structure without implying continuity between categories.

1.2 Rationale for Expanded Exploratory Analysis

The expanded EDA presented in Task 2 is intentionally broader in scope than in Task 1. This extension responds directly to feedback requesting analysis of non-diabetes outcomes and greater utilisation of the dataset. As a result, Section 3 systematically examines cardiovascular indicators, health status variables, and behavioural correlates using multiple complementary visualisations.

Given the size and multidimensional nature of the dataset, this comprehensive exploration necessarily increases the length of the EDA section. However, the extended analysis is structured purposefully: each block serves to establish demographic gradients, socioeconomic patterning, behavioural clustering, and outcome progression. The additional length reflects analytical depth rather than redundancy, and it provides the empirical basis for the revised and diversified research questions developed in Task 2.

---

2. Data Preparation

The dataset and all variable recoding steps are identical to those described in Task 1 and are reproduced in the setup chunk below (Hidden) to ensure a fully self-contained, reproducible document. No additional transformations are required for the EDA in Section 3. All 253,680 records are retained; no missing values were present.

---

3. Additional Exploratory Data Analysis: Non-Diabetes Outcomes

The following analyses explore relationships among health, lifestyle, and socioeconomic variables *independently of diabetes status*. The goal is to characterise the joint structure of these variables, in particular to understand what shapes self-reported general health, mental health burden, and functional limitation, and to motivate the non-diabetes research questions introduced in Section 4.

3.1 Cardiovascular Risk Indicators: Demographic, Socioeconomic and Behavioural Gradients

This section explores the population structure and correlates of key cardiovascular risk indicators, namely high blood pressure, high cholesterol, and stroke. The analysis moves beyond diabetes to examine how these outcomes are patterned across age, sex, socioeconomic position, and behavioural factors. Emphasis is placed on identifying demographic concentration, socioeconomic gradients, and clustering with related health behaviours. The objective is to establish structural regularities in the data that may inform subsequent modelling and risk interpretation.

3.1.1 Age–Sex Distribution by High Blood Pressure Status

```
ggplot(bp_pyramid,
       aes(x = Age, y = percent_signed, fill = Sex)) +
  geom_col(width = 0.9, colour = "white") +
  coord_flip() +
  facet_wrap(~ HighBP, nrow = 1) +
  scale_y_continuous(
    labels = function(x) paste0(abs(x), "%")
  ) +
  scale_fill_manual(
    values = pal_sex
  ) +
  labs(
    title = "Age-sex distribution by high blood pressure status",
    x     = "Age group",
    y     = "Percent within blood pressure group",
    fill  = "Sex"
  )
```

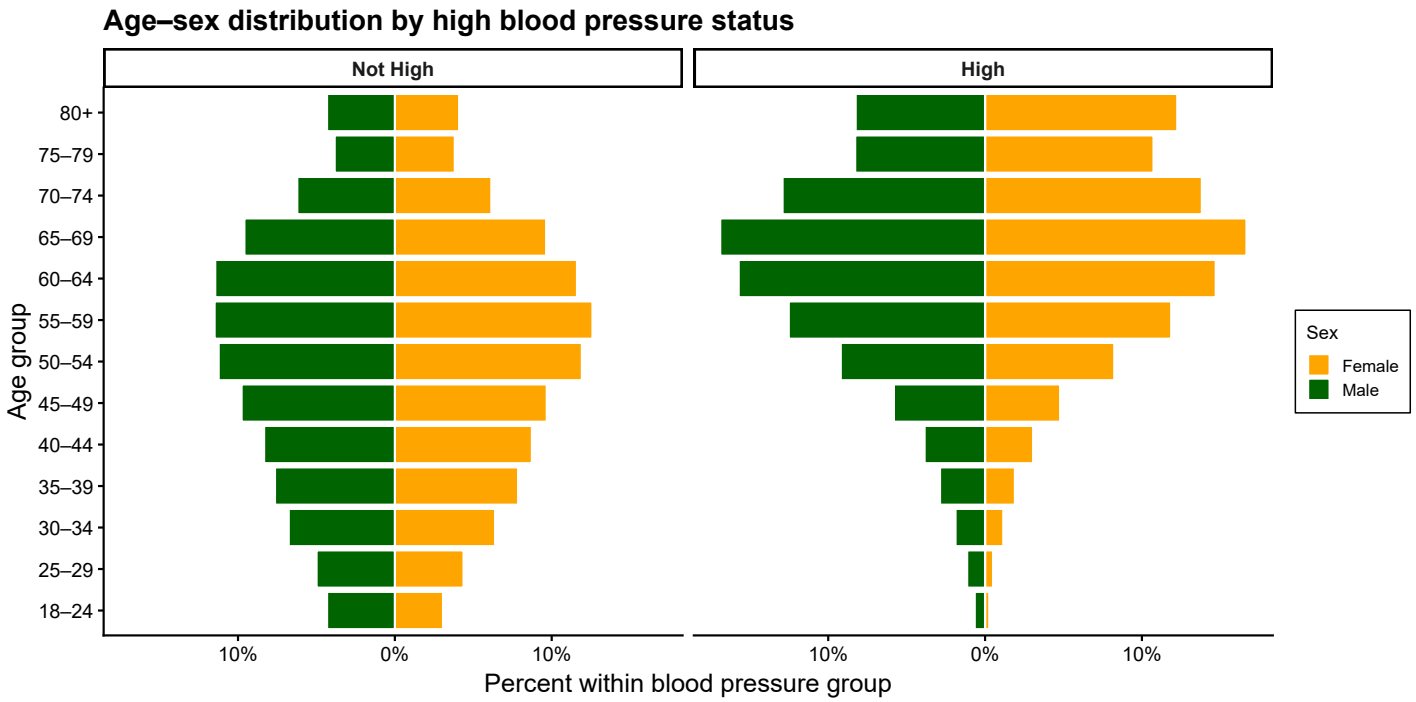

3.1.2 Age–sex distribution by high cholesterol status

```
ggplot(chol_pyramid,
       aes(x = Age, y = percent_signed, fill = Sex)) +
  geom_col(width = 0.9, colour = "white") +
  coord_flip() +
  facet_wrap(~ HighChol, nrow = 1) +
  scale_y_continuous(
    labels = function(x) paste0(abs(x), "%")
  ) +
  scale_fill_manual(
    values = pal_sex
  ) +
  labs(
    title = "Age-sex distribution by high cholesterol status",
    x     = "Age group",
    y     = "Percent within cholesterol group",
    fill  = "Sex")
```

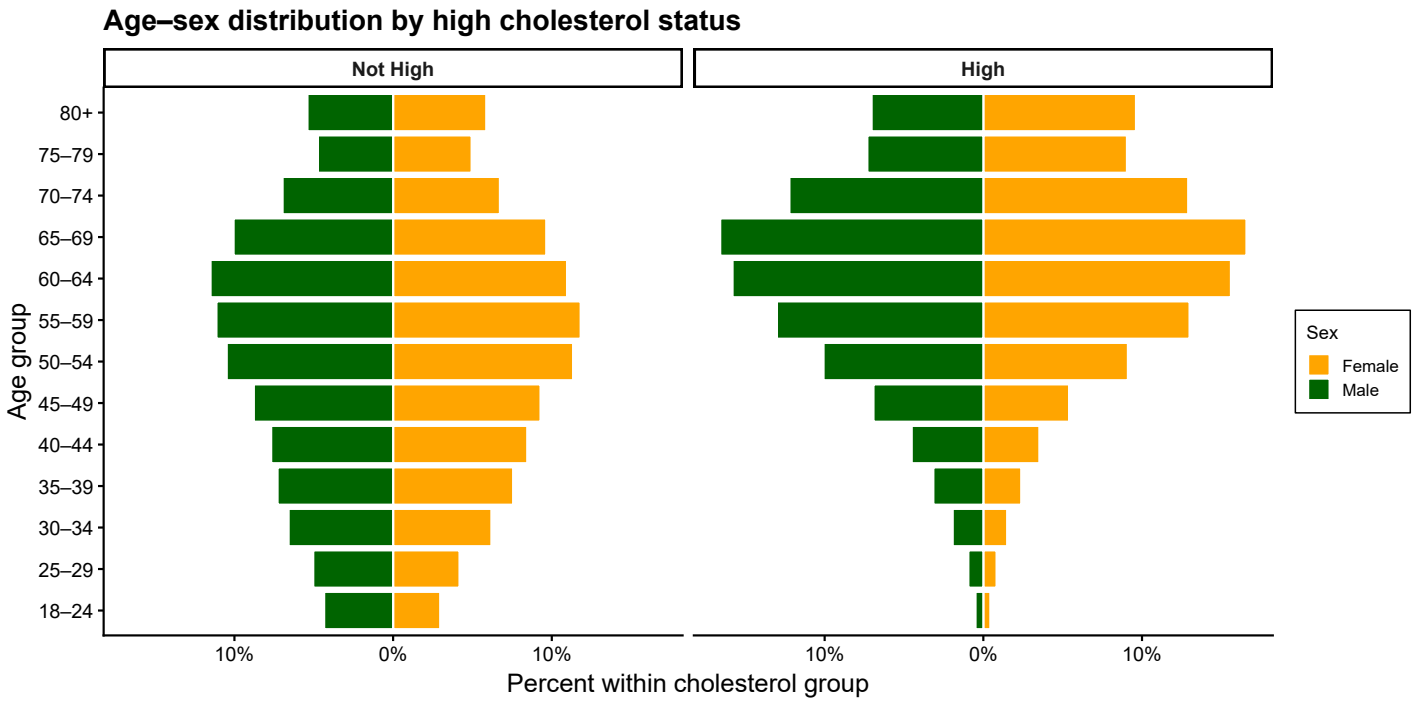

3.1.3 Age–sex distribution by stroke history

```
stroke_pyramid <- diab |>
  count(Stroke, Age, Sex) |>
  group_by(Stroke, Sex) |>
  mutate(
    percent = 100 * n / sum(n)
  ) |>
  ungroup() |>
  mutate(
    percent_signed = case_when(
      Sex == "Male"   ~ -percent,
      Sex == "Female" ~  percent,
      TRUE            ~ NA_real_
    )
  )

ggplot(stroke_pyramid,
       aes(x = Age, y = percent_signed, fill = Sex)) +
  geom_col(width = 0.9, colour = "white") +
  coord_flip() +
  facet_wrap(~ Stroke, nrow = 1) +
  scale_y_continuous(
    labels = function(x) paste0(abs(x), "%")
  ) +
  scale_fill_manual(
    values = pal_sex
  ) +
  labs(
    title = "Age-sex distribution by stroke history",
    x     = "Age group",
    y     = "Percent within stroke-history group",
    fill  = "Sex"
```

```
  )
```

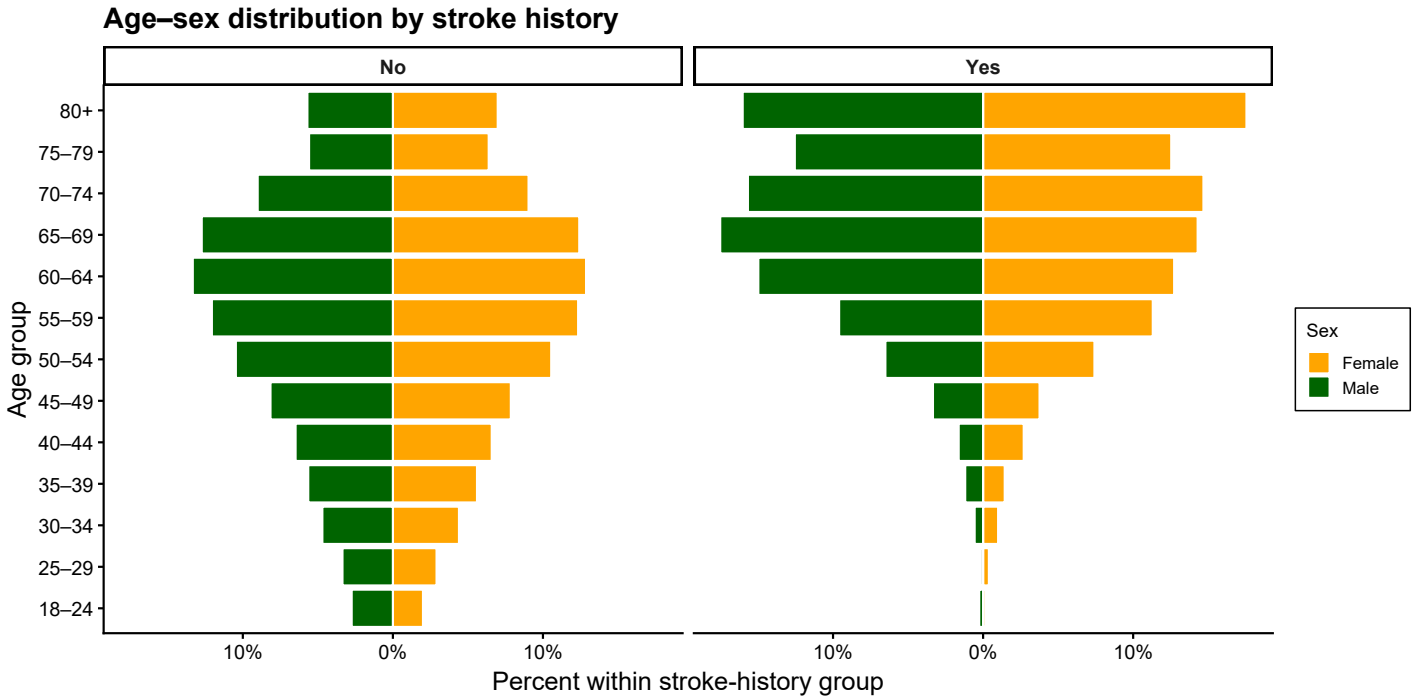

Across all three outcomes, high blood pressure, high cholesterol, and stroke, the dominant structural feature is age concentration. Each "Yes" panel is heavily skewed toward the 55–69 and older age bands, while representation below age 40 is minimal. This confirms that cardiovascular burden in this dataset is fundamentally age-structured rather than evenly distributed across the life course.

However, the degree of age concentration differs meaningfully between outcomes. High blood pressure and high cholesterol show similar mid-to-late life clustering, with substantial presence beginning in the 50–54 band and peaking around 60–69. Stroke, by contrast, exhibits a far sharper right-skew: the "Yes" group is overwhelmingly concentrated in older age bands, with negligible representation in younger adults. This suggests stroke is not merely age-associated but strongly age-dependent relative to the other indicators.

Sex differences are comparatively modest within age strata for high blood pressure and high cholesterol. Once age is conditioned upon, male–female distributions appear broadly parallel, implying that sex operates as a secondary structuring factor relative to age. In the stroke group, however, there is a slight female predominance in the oldest age bands, which may reflect differential survival into advanced age rather than a pure sex-specific risk gradient.

Importantly, these pyramids depict the **composition within outcome groups**, not prevalence across the population. They therefore highlight demographic concentration rather than risk probability. The consistency across high blood pressure and high cholesterol suggests shared cardiometabolic ageing processes, whereas stroke displays a more extreme demographic compression, indicating it may represent the culmination of accumulated vascular risk rather than an early-stage cardiovascular condition.

Taken together, the figures establish age as the primary axis of cardiovascular risk patterning in this sample. Any socioeconomic or behavioural gradients observed later must therefore be interpreted within the context of this strong underlying demographic structure.

3.1.4 Joint SES heatmap of High Blood Presure prevalence.

```
bp_ses |>
  ggplot(aes(x = Income, y = Education, fill = prev_bp)) +
  geom_tile(colour = "white", linewidth = 1.2) +
  geom_text(
```

```
    aes(label = percent(prev_bp, accuracy = 1)),
    colour   = "white",
    fontface = "bold",
    size     = 4.5) +
  scale_fill_gradient(
    low  = "#F1EEF6",
    high = "#882255",
    name = "Prevalence\nHigh BP",
    labels = percent_format(accuracy = 1)) +
  labs(title = "Prevalence of high blood pressure by education and income",
    x     = "Income band",
    y     = "Education level")
```

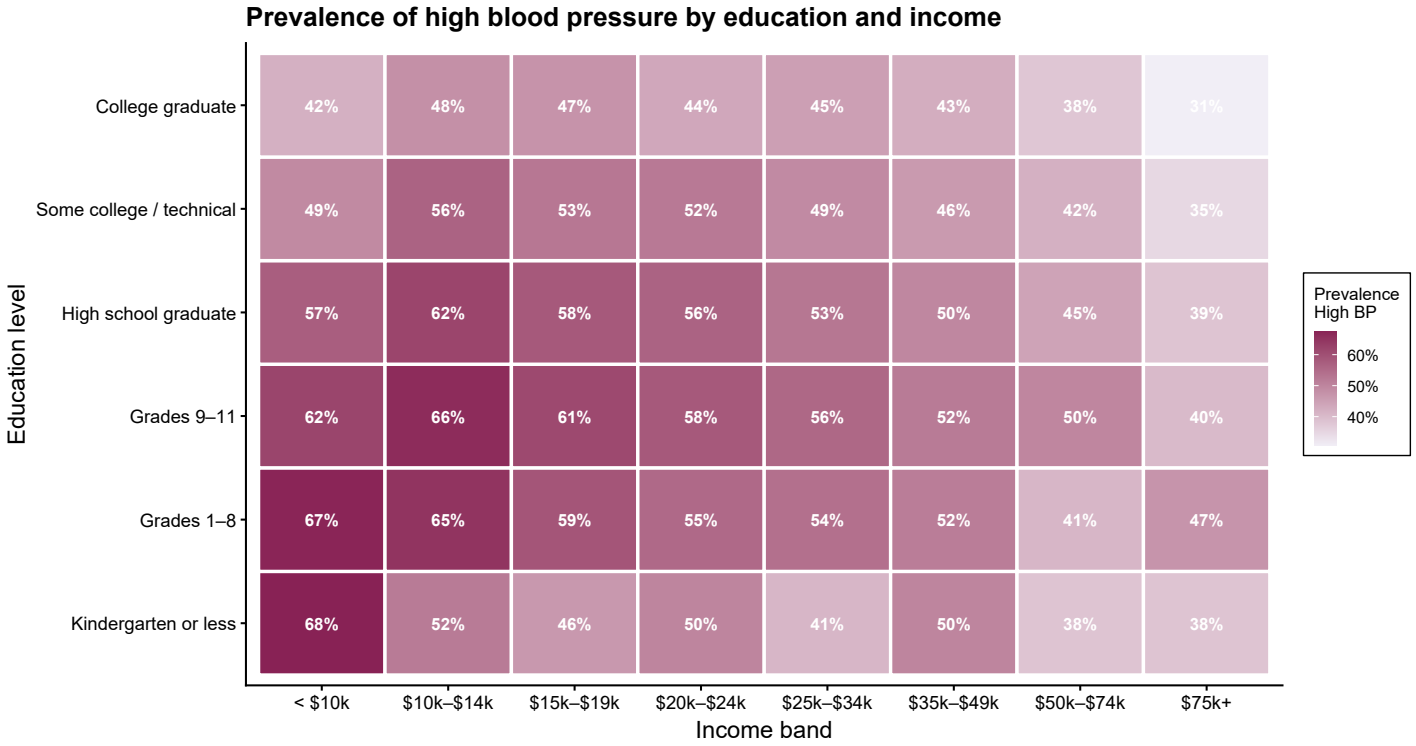

3.1.5 Joint SES heatmap of High Cholestrol prevalence.

```
chol_ses |>
  ggplot(aes(x = Income, y = Education, fill = prev_chol)) +
  geom_tile(colour = "white", linewidth = 1.2) +
  geom_text(
    aes(label = percent(prev_chol, accuracy = 1)),
    colour   = "white",
    fontface = "bold",
    size     = 3.5
  ) +
  scale_fill_gradient(
    low    = "#F1EEF6",
    high   = "#882255",
    name   = "Prevalence\nHigh cholesterol",
    labels = percent_format(accuracy = 1)
  ) +
  labs(
    title = "Prevalence of high cholesterol by education and income",
```

```
    x      = "Income band",
    y      = "Education level"
  )
```

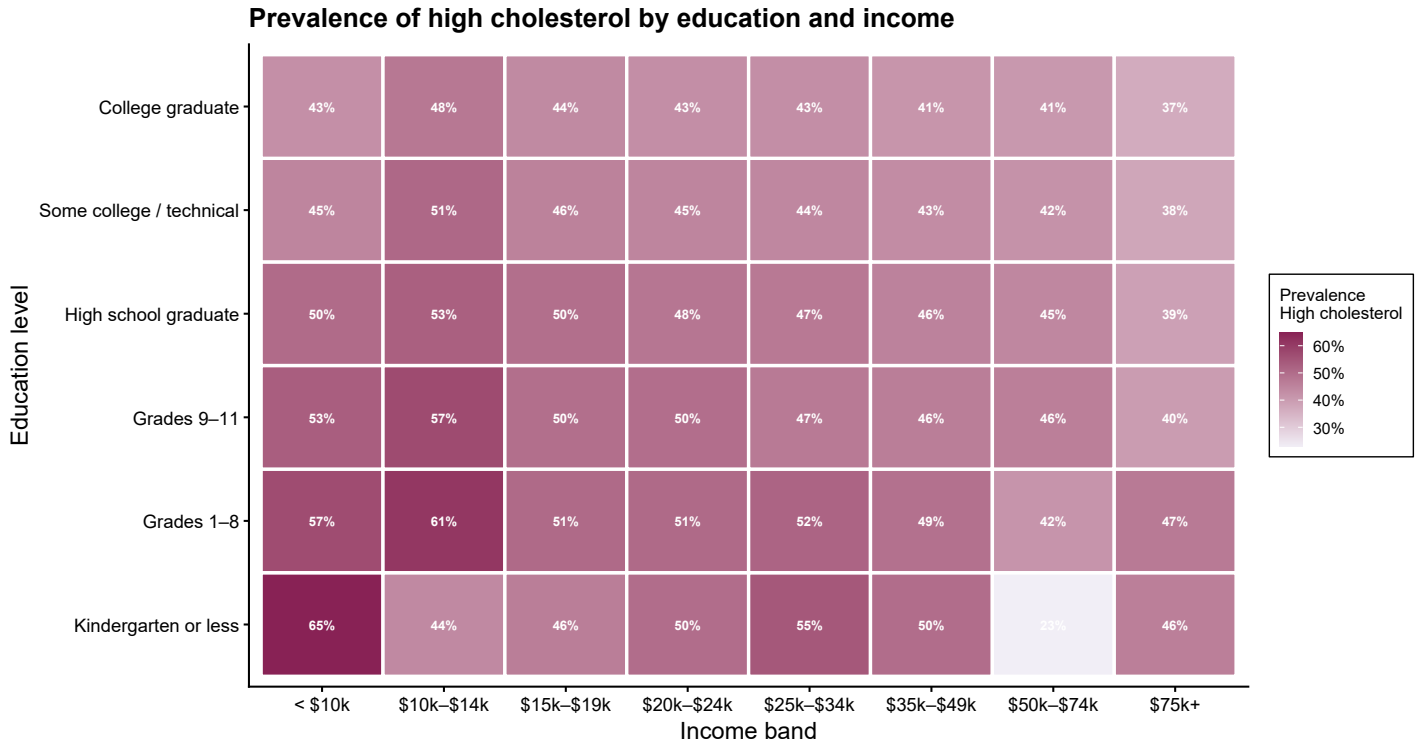

3.1.6 Joint SES heatmap of Stroke prevalence.

```
stroke_ses |>
  ggplot(aes(x = Income, y = Education, fill = prev_stroke)) +
  geom_tile(colour = "white", linewidth = 1.2) +
  geom_text(
    aes(label = percent(prev_stroke, accuracy = 1)),
    colour = "white",
    fontface = "bold",
    size = 3.5) +
  scale_fill_gradient(
    low = "#F1EEF6",
    high = "#882255",
    name = "Prevalence\nStroke history",
    labels = percent_format(accuracy = 1)) +
  labs(title = "Prevalence of stroke history by education and income",
       x = "Income band",
       y = "Education level")
```

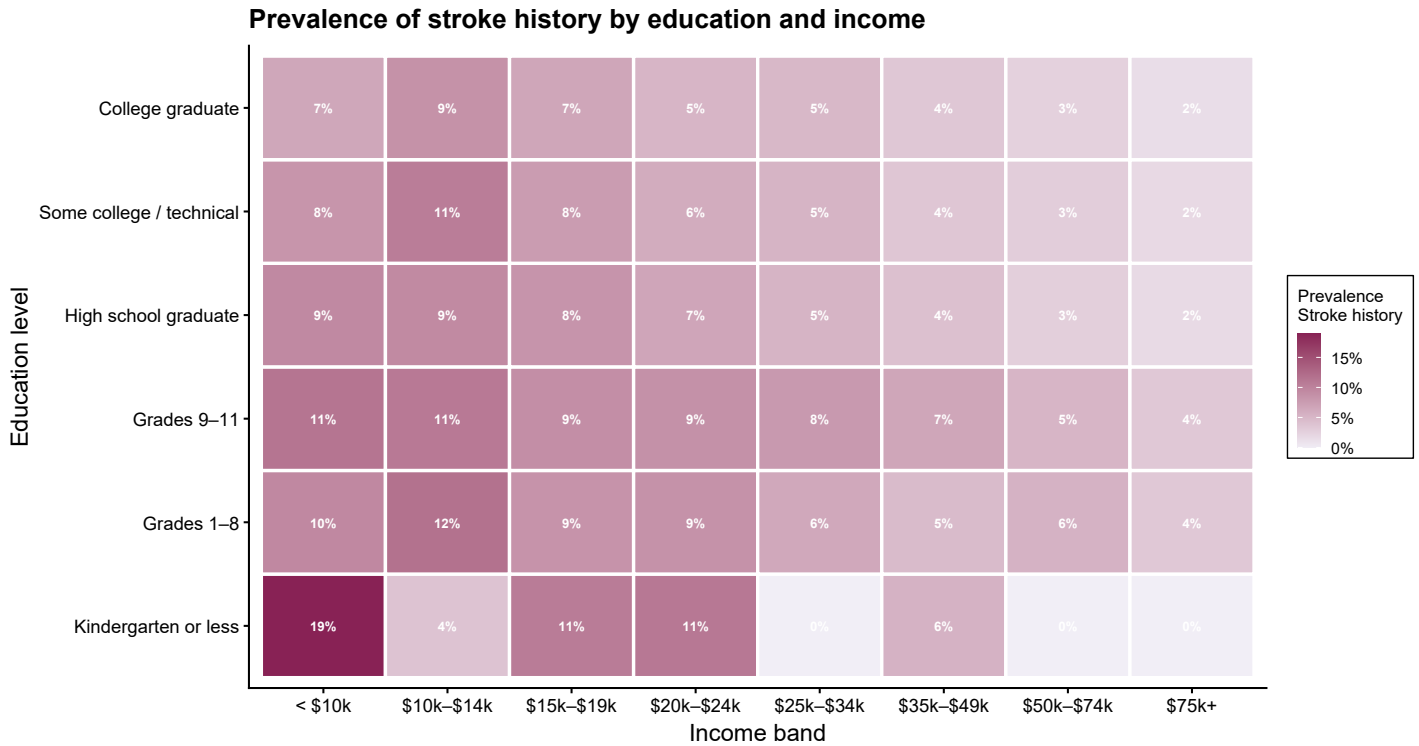

Across all three outcomes, a clear socioeconomic gradient is evident. Prevalence is consistently highest in the lowest education and lowest income strata and declines with increasing educational attainment and income. The lower-left cells (least educated, lowest income) exhibit the greatest burden, while college graduates in the highest income bands show the lowest prevalence. This pattern indicates cumulative socioeconomic disadvantage rather than a single-dimension effect.

High blood pressure and high cholesterol display particularly strong and stable two-dimensional gradients. Within nearly every education category, prevalence decreases progressively as income rises. Likewise, within each income band, higher education corresponds to lower prevalence. The monotonic pattern across both axes suggests that both income and education independently structure cardiometabolic risk, rather than one merely proxying for the other.

Stroke exhibits the same directional gradient but at substantially lower absolute prevalence and with more variability across cells. Some low-education, low-income cells show markedly elevated stroke prevalence (e.g., 19%), while higher SES cells approach very low levels. The noisier pattern likely reflects smaller counts and the relative rarity of stroke, but the overall concentration in socioeconomically disadvantaged groups remains clear.

Importantly, these are crude prevalence estimates. As shown in the age–sex pyramids, lower socioeconomic strata contain disproportionately older individuals, and age is a dominant determinant of all three outcomes. Therefore, part of the observed SES gradient likely reflects demographic composition. However, the persistence of a consistent gradient across both income and education dimensions suggests that socioeconomic position is not merely a proxy for age but a meaningful structural correlate of cardiovascular burden.

3.1.7 Prevalence Cardiovascular Outcomes Accumulated Risk Factors

```
ggplot(cvd_long, aes(x = risk_group, y = prop, fill = risk_group)) +
  geom_col(width = 0.7) +
  facet_wrap(~ Outcome, nrow = 1) +
  scale_y_continuous(
    labels = percent_format(accuracy = 1),
    limits = c(0, 1)
  ) +
  scale_fill_brewer(palette = "Purples", name = "Risk\ncount") +
```

```
  labs(
    title = "Prevalence of cardiovascular outcomes by cardiometabolic risk factor count",
    x     = "Number of risk factors present\n(obesity, inactivity, smoking, diabetes)",
    y     = "Proportion with outcome"
  )
```

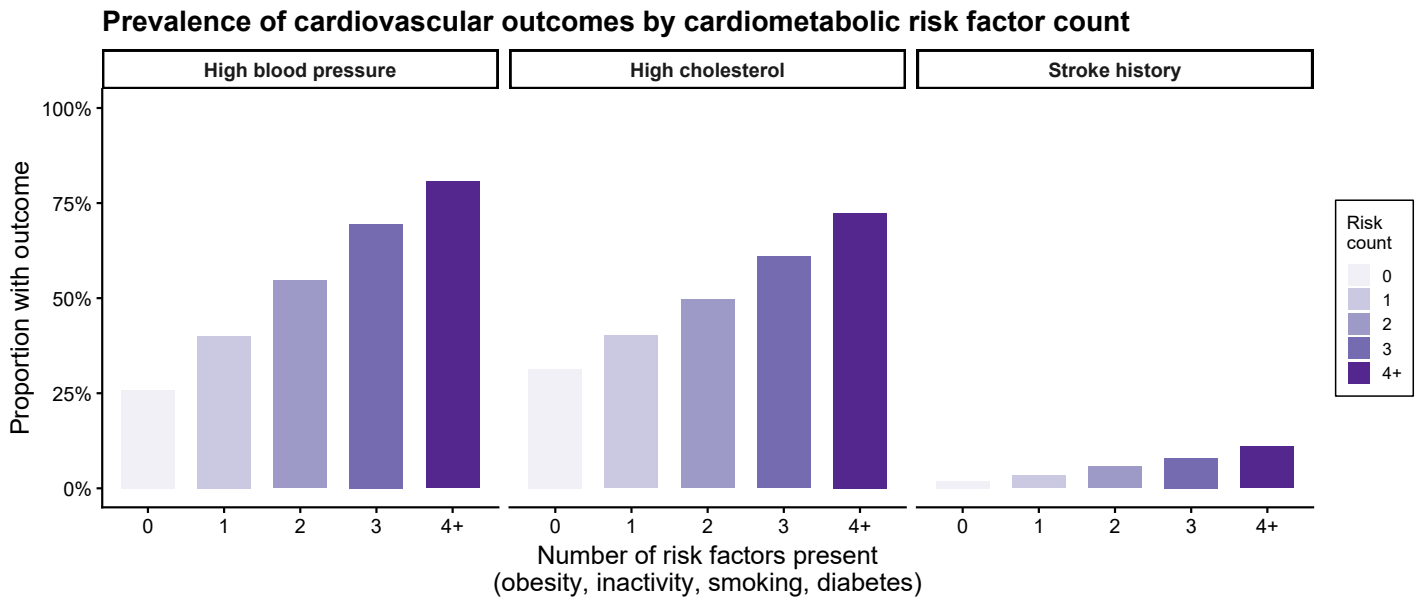

Cardiovascular outcomes show a clear **dose–response relationship with the accumulation of cardiometabolic risk factors**. For both **high blood pressure** and **high cholesterol**, prevalence increases steadily as additional risk factors are present. Individuals with no risk factors show relatively low prevalence (around one quarter to one third), while those with **four or more risk factors exhibit very high prevalence**, exceeding roughly **70–80%** for both conditions. This pattern indicates that these outcomes are strongly associated with the combined burden of obesity, physical inactivity, smoking, and diabetes.

Stroke history follows the same upward trend, but at substantially lower levels overall. Even among individuals with multiple risk factors, stroke prevalence remains below roughly **10–12%**, reflecting the relative rarity of stroke compared with hypertension and hypercholesterolaemia in the population.

Overall, the figure demonstrates that **cardiometabolic risk factors act cumulatively rather than independently**. As risk exposures accumulate, the likelihood of cardiovascular conditions, particularly hypertension and high cholesterol, increases markedly, highlighting the importance of addressing multiple lifestyle and metabolic risk factors simultaneously.

3.1.8 Joint blood pressure and cholesterol profiles across age groups

```
ggplot(bp_chol_age,
       aes(x = Age, y = prop, fill = interaction(HighBP, HighChol,
                                                 sep = " / Chol: "))) +
  geom_col(position = "fill", width = 0.85) +
  scale_y_continuous(labels = percent_format(accuracy = 1)) +
  scale_fill_viridis_d(option = "D", direction = -1) +
  labs(
    title = "Joint blood pressure and cholesterol status by age group",
    x     = "Age group",
    y     = "Proportion within age group",
    fill  = "BP / Cholesterol"
  )
```

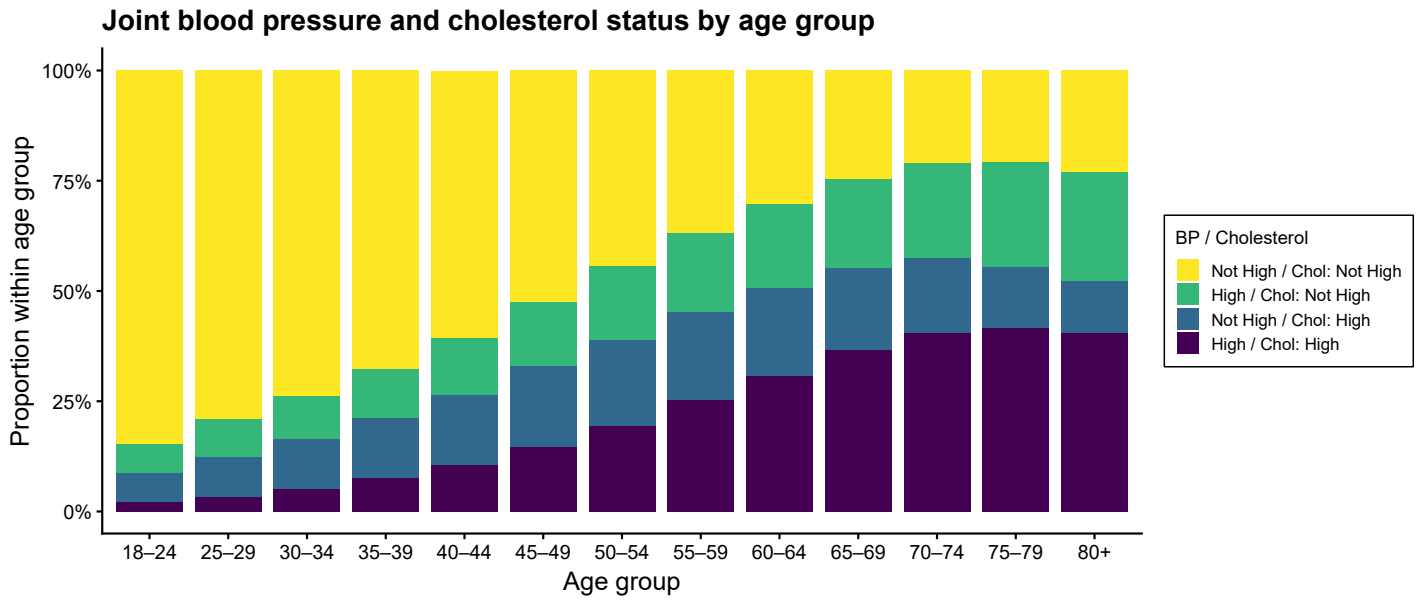

The distribution of cardiovascular profiles shifts systematically with age. Younger groups are overwhelmingly characterised by normal blood pressure and cholesterol, whereas the proportion with dual elevation (high/high) increases sharply from midlife onward and becomes dominant in older age bands. This pattern reflects progressive cardiometabolic accumulation rather than isolated risk emergence.

Importantly, the high/high category expands more rapidly than either single-risk category, indicating clustering rather than independence of these conditions. By ages 65+, dual dysregulation accounts for a substantial share of the population within each age group, suggesting convergence of vascular risk factors with ageing.

The pattern implies that hypertension and hypercholesterolaemia are not parallel but interacting processes within cardiovascular ageing, reinforcing their role as intermediate markers in the pathway toward more severe vascular outcomes.

3.1.9 Co-occurrence of Hypertension and High Cholesterol

```
diab |>
  count(HighBP, HighChol) |>
  mutate(prop = n / sum(n)) |>
  ggplot(aes(x = HighBP, y = HighChol, fill = prop)) +
  geom_tile(colour = "white", linewidth = 1.5) +
  geom_text(
    aes(label = percent(prop,accuracy = 0.1)),
    size     = 6,
    fontface = "bold",
    colour   = "white"
  ) +
  labs(
    title = "Joint prevalence of high blood pressure and high cholesterol",
    x     = "High blood pressure",
    y     = "High cholesterol"
  ) +
  theme(
    legend.position = 'none'
  )
```

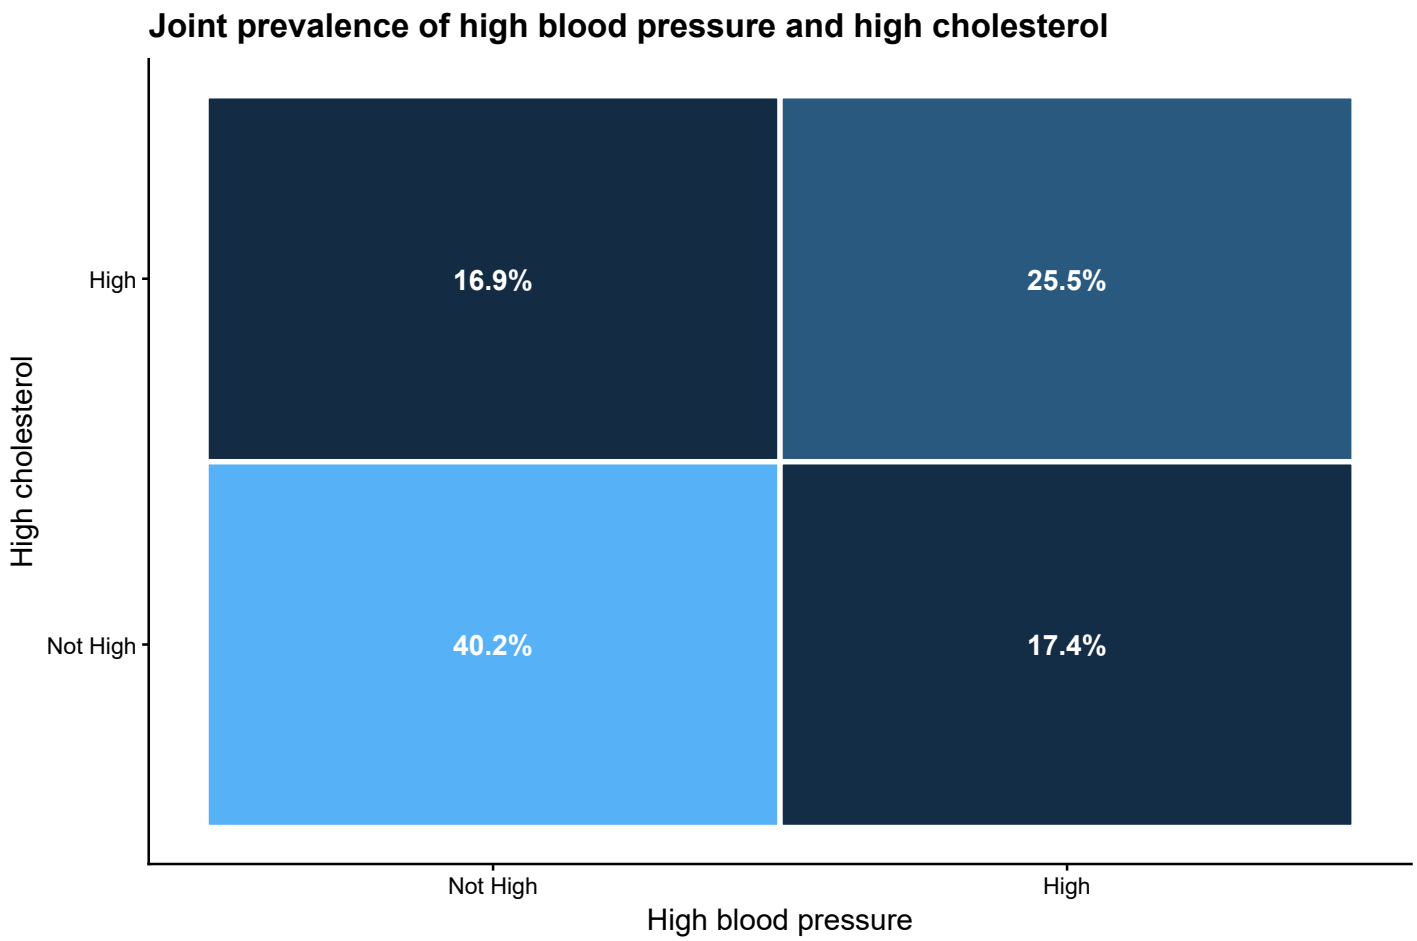

Approximately **40.2%** of the sample has neither high blood pressure nor high cholesterol. The remaining **59.8%** have at least one of the two conditions, indicating that cardiometabolic risk is common in this population.

The largest single elevated profile is **concurrent high blood pressure and high cholesterol (25.5%)**, which exceeds either isolated condition (17.4% BP only; 16.9% cholesterol only). This suggests meaningful clustering rather than independent occurrence of the two risk factors.

Overall, hypertension and hypercholesterolaemia appear strongly co-patterned, with dual elevation more common than either single-condition state, reinforcing their joint role in cardiovascular risk stratification.

3.1.10 Stroke prevalence by blood pressure and cholesterol status

```
ggplot(stroke_bpchol,
       aes(x = BP_Chol, y = prop, fill = prop)) +
  geom_col(width = 0.55) +
  geom_text(aes(label = percent(prop, accuracy = 0.1)),
            vjust = -0.4, size = 5.5, fontface = "bold") +
  scale_y_continuous(labels = percent_format(accuracy = 0.1),
                     limits = c(0, 0.12)) +
  scale_fill_viridis_c(option = "C", direction = -1, guide = "none") +
  labs(
    title = "Stroke prevalence by blood pressure and cholesterol status",
    x     = "Blood pressure / cholesterol profile",
    y     = "Proportion with stroke history"
  )
```

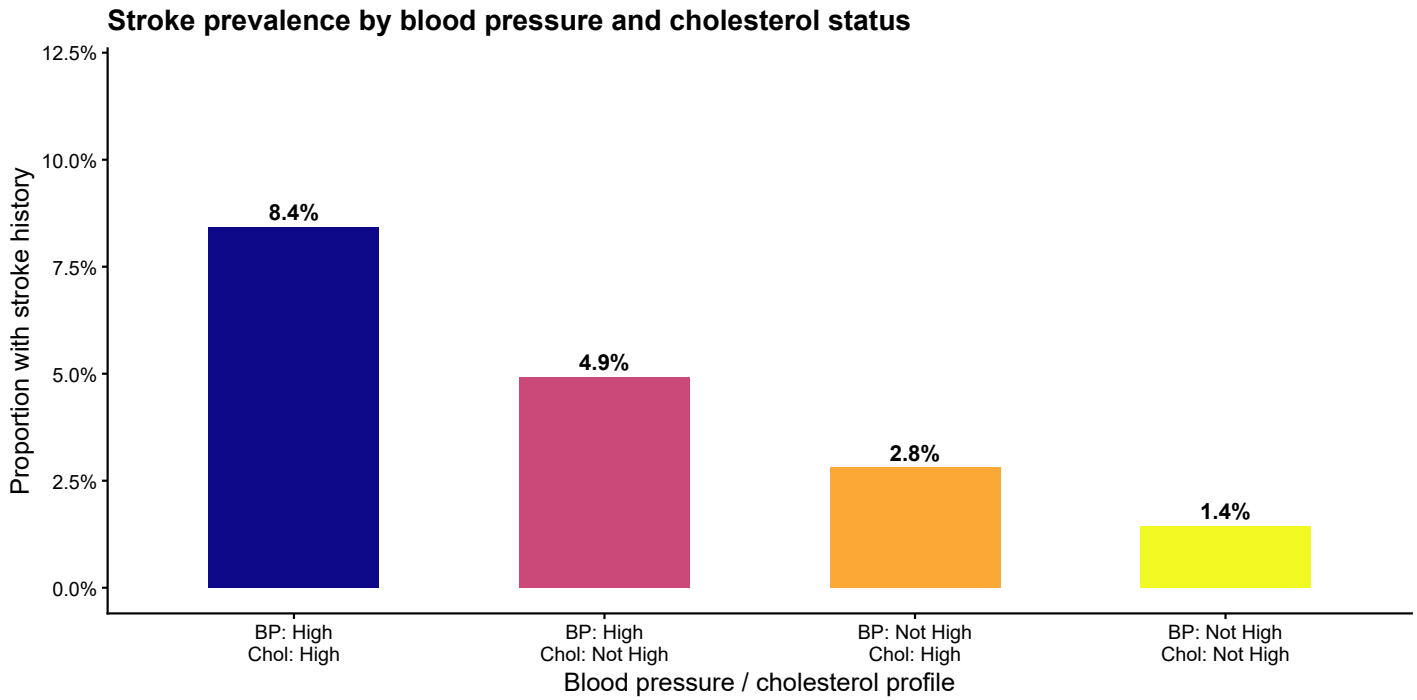

Stroke prevalence increases monotonically across cardiovascular profiles. The highest prevalence is observed among individuals with both high blood pressure and high cholesterol, followed by isolated high blood pressure, isolated high cholesterol, and neither condition.

The gradient suggests cumulative vascular burden, with blood pressure appearing more strongly associated with stroke than cholesterol alone. Stroke thus reflects downstream risk concentration rather than isolated cardiometabolic dysfunction.

3.1.11 Joint co-occurrence profiles of high blood pressure, high cholesterol, and stroke

```
ggplot(joint, aes(x = reorder(Profile, Proportion), y = Proportion, fill = Proportion)) +
  geom_col(width = 0.65) +
  geom_text(
    aes(label = percent(Proportion, accuracy = 0.1)),
    hjust = -0.15,
    size  = 4.5
  ) +
  coord_flip() +
  scale_x_discrete(expand = expansion(mult = c(0.05, 0.05))) +
  scale_y_continuous(
    labels = percent_format(accuracy = 1),
    limits = c(0, max_prop * 1.15)
  ) +
  scale_fill_viridis_c(option = "D", direction = -1, guide = "none") +
  labs(
    title = "Co-occurrence profiles for HighBP, HighChol, and Stroke",
    x     = NULL,
    y     = "Proportion of full sample"
  )
```

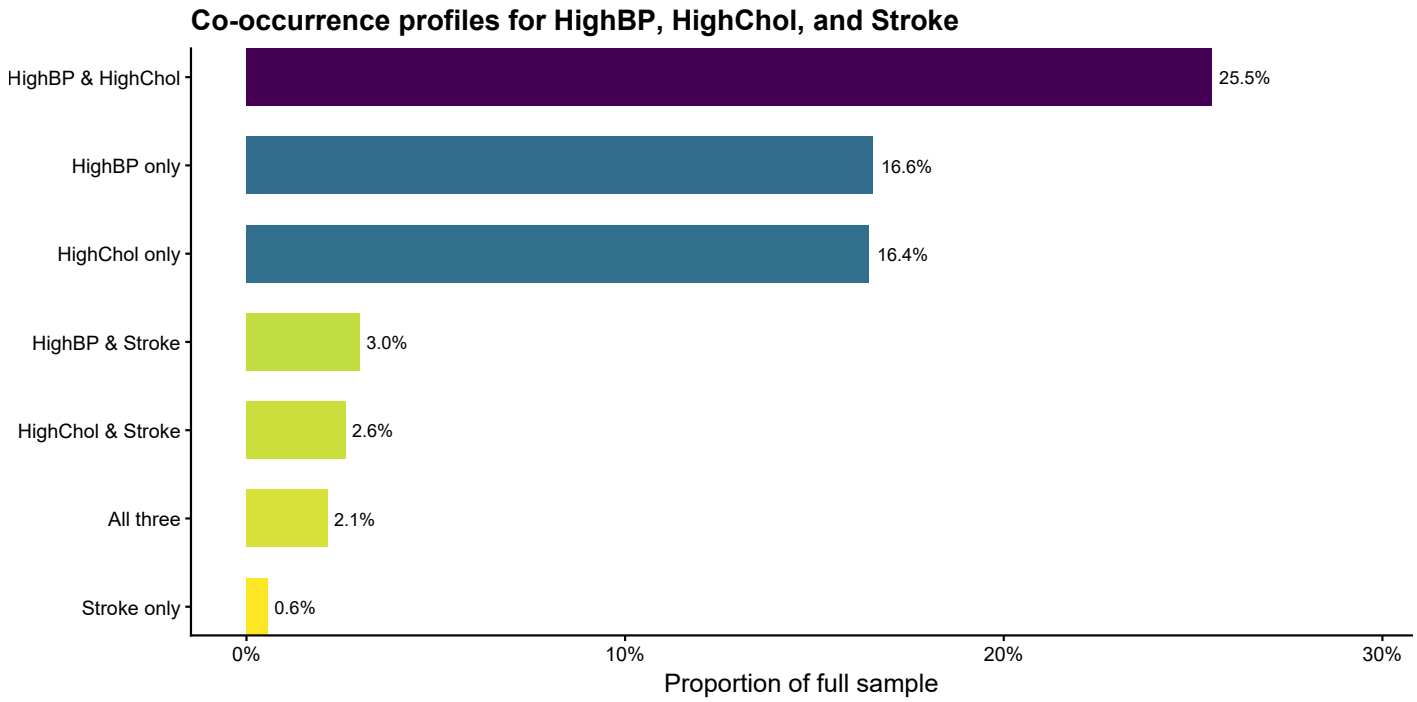

3.1.12 Association Strength - Cramér's V

```
ggplot(cv_nondm,
       aes(x = reorder(Predictor, CramersV),
           y = CramersV,
           fill = Outcome)) +
  geom_col(position = position_dodge(width = 0.75), width = 0.7) +
  geom_text(aes(label = round(CramersV, 2)),
            position = position_dodge(width = 0.75),
            hjust = -0.1, size = 3.2) +
  coord_flip() +
  expand_limits(x = max(cv_nondm$CramersV) + 0.03) +
  scale_fill_manual(values = c(
    "High Blood Pressure" = "#2c7bb6",
    "High Cholesterol"    = "#f46d43",
    "Stroke"              = "#1a9641"
  )) +
  scale_y_continuous(limits = c(0, 0.45)) +
  labs(
    title = "Association strength with non-diabetes outcomes (Cramér's V)",
    x     = NULL,
    y     = "Cramér's V",
    fill  = "Outcome"
  ) +
  theme(legend.position = "bottom")
```

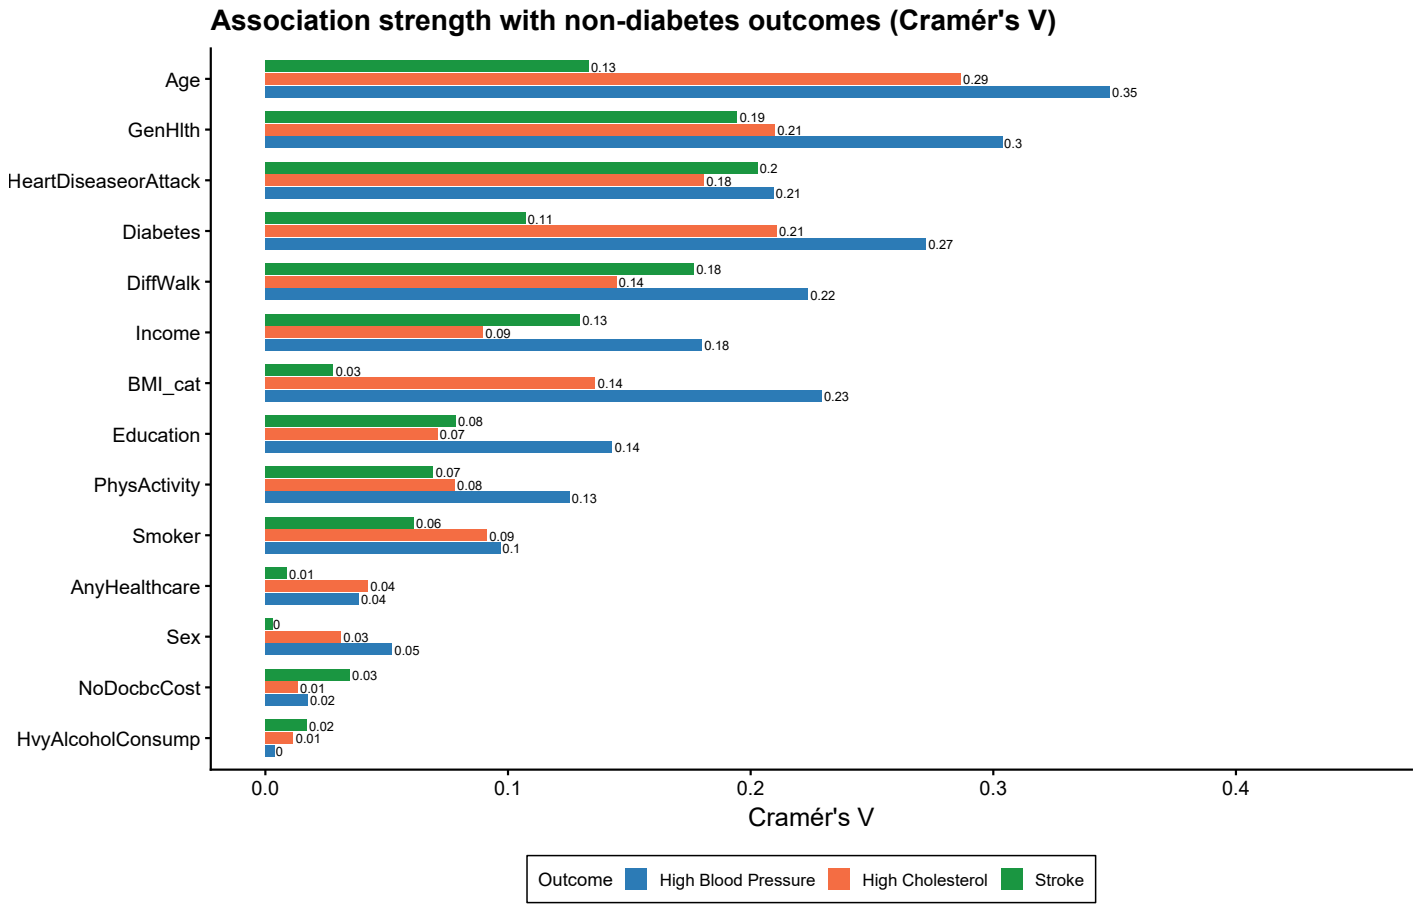

This chart highlights how different risk factors weigh against cardiovascular outcomes using **Cramér's V** (a measure of association where higher values indicate a stronger relationship).

The data reveals two distinct "profiles" of disease:

1. **Metabolic & Age-Driven (High BP & Cholesterol)**

Hypertension and High Cholesterol are primarily driven by **systemic, long-term factors**.

- **Primary Drivers: Age** ($V \approx 0.35$) and **General Health** ($V \approx 0.30$) are the strongest predictors.
- **Secondary Drivers:** Strong links to **Diabetes** and **BMI**, suggesting these are markers of general metabolic decline.
- **Takeaway:** These conditions follow a broad, population-wide risk gradient shaped by getting older and overall physical state.

2. **Clinical & Pathological (Stroke)**

Stroke follows a different logic, focusing more on **existing damage** than just age.

- **Primary Drivers: Prior Heart Disease** ($V \approx 0.20$), **General Health** ($V \approx 0.19$), and **Walking Difficulty** ($V \approx 0.18$).
- **Key Difference:** Age is significantly less dominant here than it is for Blood Pressure. BMI has almost no direct association with Stroke in this specific view.
- **Takeaway:** Stroke appears as a "clinical endpoint", it is most strongly associated with patients who already show signs of cardiovascular failure or physical impairment.

Outcome	Strongest Predictor	Second Predictor
High Blood Pressure	Age (0.35)	General Health (0.30)
High Cholesterol	Age (0.29)	General Health (0.21)

Stroke	Heart Disease (0.20)	General Health (0.19)

3.1.13 Summary of Key Findings

Section 3.1 establishes that cardiovascular risk in this dataset is structured along three primary axes: **age, socioeconomic position, and cardiometabolic clustering**.

First, **age is the dominant organising factor**. High blood pressure and high cholesterol rise steadily from midlife onward, while stroke is far more sharply concentrated in older age groups. Stroke appears less as a gradual condition and more as a late-stage vascular outcome.

Second, a **clear and consistent socioeconomic gradient** is evident across all three outcomes. Lower education and lower income strata show substantially higher crude prevalence, with monotonic declines across both dimensions. Although partly influenced by age composition, the two-dimensional gradient suggests an independent structural role for socioeconomic position.

Third, **cardiometabolic conditions cluster rather than occur independently**. Dual high blood pressure and high cholesterol (25.5%) exceed either isolated condition, and stroke prevalence increases monotonically across BP/cholesterol profiles, peaking in the dual-high group. This supports a cumulative vascular burden framework.

Fourth, **obesity strongly structures hypertension and cholesterol but not stroke**, reinforcing the interpretation of BP and cholesterol as intermediate metabolic states. Stroke instead aligns more strongly with prior heart disease, general health, and walking difficulty.

Finally, association analyses (Cramér's V) reveal two distinct disease profiles:

- **Hypertension and cholesterol** are predominantly age and general-health driven.
- **Stroke** is more closely associated with established cardiovascular pathology and functional impairment.

Overall, Section 3.1 demonstrates that cardiovascular outcomes in this dataset are not randomly distributed but follow coherent demographic, socioeconomic, and pathological structures. These findings provide the empirical foundation for the modelling strategy developed in Section 4.

3.1.14 Implications for Modelling

The exploratory analysis of cardiometabolic outcomes suggests several considerations for subsequent modelling. First, the strong accumulation pattern observed across cardiometabolic risk factors indicates that lifestyle and metabolic variables such as obesity, smoking, physical inactivity, and diabetes are likely to act as key predictors of cardiovascular outcomes. Second, the association strength analysis shows that demographic and health-status variables—particularly age and self-rated general health—have the strongest relationships with high blood pressure and high cholesterol, indicating that these variables should be prioritised in predictive models.

For stroke, the strongest associations appear with prior cardiovascular disease history, poorer general health, and functional limitation (difficulty walking), suggesting that stroke risk may be more closely linked to underlying health status and comorbidity than to behavioural factors alone. Finally, the relatively low prevalence of stroke compared with hypertension and hypercholesterolaemia suggests that models predicting stroke may face class imbalance and may require approaches suitable for rare outcomes.

Overall, the EDA indicates that **demographic factors, general health status, and cardiometabolic risk profiles** are likely to be central components of the modelling framework used to investigate cardiovascular outcomes.

3.2 Patterns and determinants of mental health burden

This section examines the distribution and key correlates of **mental health burden**, measured as the number of mentally unhealthy days reported in the past 30 days. The analysis explores how mental health burden varies across demographic, socioeconomic, and health-related factors, including age, physical health, and lifestyle behaviours. By identifying systematic patterns in mental health outcomes, this section aims to highlight the factors most strongly associated with elevated psychological distress in the population.

3.2.1 Distribution of mentally unhealthy days in the past 30 days

```
ggplot(diab, aes(x = MentHlth)) +
  geom_histogram(
    binwidth = 1,
    boundary = 0,
    fill = "steelblue",
    colour = "black"
  ) +
  scale_x_continuous(
    breaks = seq(0, 30, by = 5),
    limits = c(0, 30)
  ) +
  scale_y_continuous(
    labels = label_number(scale = 1e-3, suffix = "K"),
    breaks = seq(0,200000,by = 25000)
  ) +
  labs(
    title = "Distribution of mentally unhealthy days in the past 30 days",
    x = "Number of mentally unhealthy days",
    y = "Frequency")
```

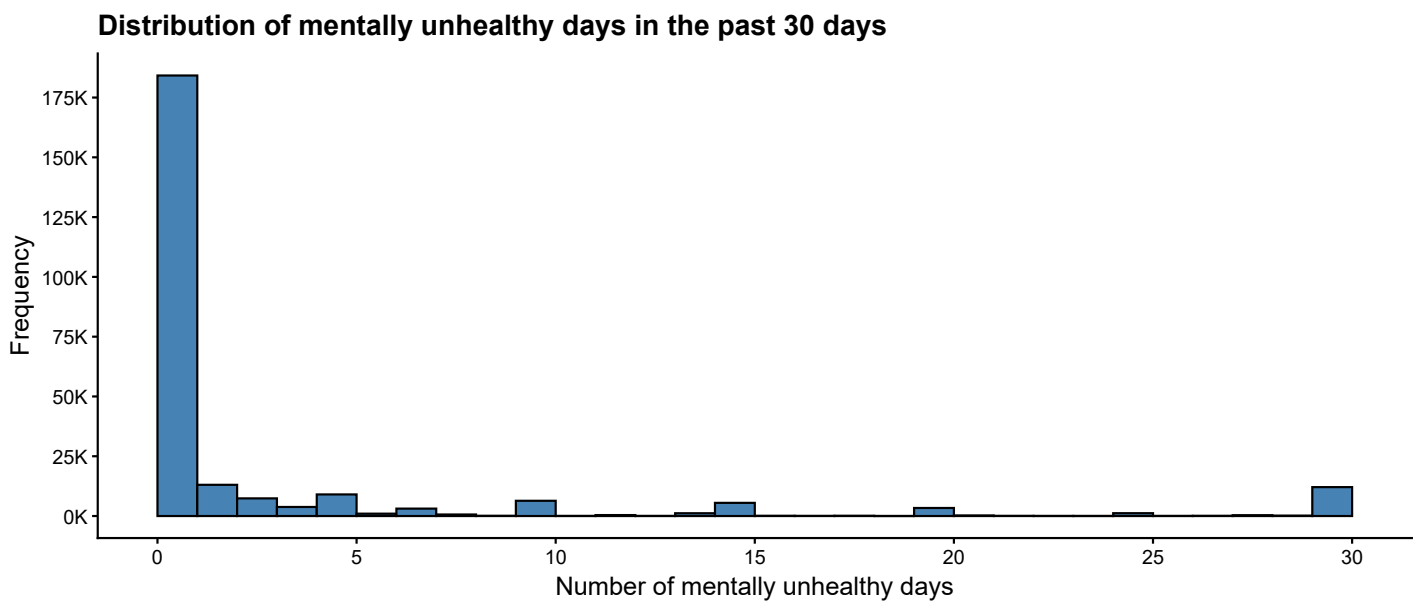

The distribution of mentally unhealthy days is **highly right-skewed**, with a very large proportion of respondents reporting **zero mentally unhealthy days in the past month**. Frequencies decline sharply as the number of unhealthy days increases, producing a long right tail that extends toward the upper bound of 30 days. This pattern indicates that while most individuals report little or no recent mental health burden, a smaller subset experiences **frequent or persistent distress**.

The distribution also shows evidence of **excess zeros and strong dispersion**, with the spike at zero far exceeding the frequencies of neighbouring values. Such a structure suggests that standard linear modelling would be inappropriate, and motivates the use of **count-based models**, potentially including **zero-inflated or overdispersed models (e.g., Poisson or zero-inflated negative binomial)** when analysing mental health burden as an outcome in subsequent modelling.

3.2.3 Mental Health Burden by Income and Sex

```r
diab |>
  mutate(HighMH = MentHlth >= 14) |>
  group_by(Income, Sex) |>
  summarise(prop_high = mean(HighMH), .groups = "drop") |>

  ggplot(aes(x = Income, y = prop_high, fill = Sex)) +
  geom_col(
    position = position_dodge(width = 0.75),
    width = 0.7,
    colour = "black",
    linewidth = 0.2
  ) +
  geom_text(
    aes(label = percent(prop_high, accuracy = 1)),
    position = position_dodge(width = 0.9),
    vjust = -0.3, size = 4) +
  scale_y_continuous(
    labels = percent_format(accuracy = 1),
    limits = c(0, NA)
  ) +
  scale_fill_manual(
    values = pal_sex
  ) +
  labs(
    title = "Prevalence of high mental health burden ( 14 days) by income and sex",
    x     = "Household income",
    y     = "Proportion with  14 mentally unhealthy days",
    fill  = "Sex"
  )
```

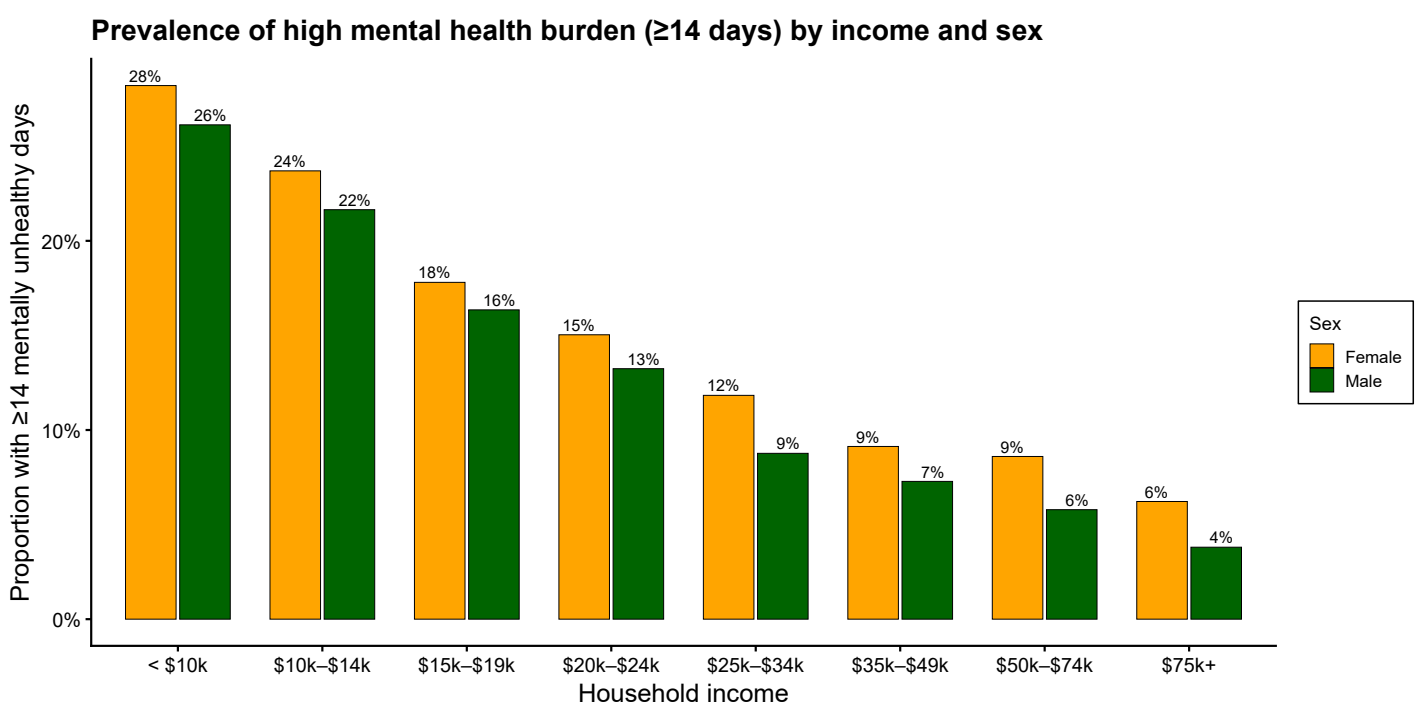

A clear socioeconomic gradient is evident in the prevalence of high mental health burden. The proportion of respondents reporting ≥14 mentally unhealthy days declines steadily as household income increases, indicating that individuals in lower-income groups experience substantially higher levels of psychological distress. In the lowest income category (<$10k), roughly 28% of females and 26% of males report high mental health burden,

compared with only 6% and 4% respectively in the highest income group ($75k+).

Across all income categories, females consistently report higher prevalence than males, although the difference between sexes is modest relative to the overall income gradient. These results suggest that socioeconomic position is a strong determinant of mental health burden, with financial disadvantage associated with markedly higher levels of frequent psychological distress. (Money doesn't buy happiness :))

3.2.3 Mental Health Burden by Educatioand Sex

```
diab |>
  mutate(HighMH = MentHlth >= 14) |>
  group_by(Education, Sex) |>
  summarise(prop_high = mean(HighMH), .groups = "drop") |>

  ggplot(aes(x = Education, y = prop_high, fill = Sex)) +
  geom_col(
    position = position_dodge(width = 0.75),
    width = 0.7,
    colour = "black",
    linewidth = 0.2
  ) +
  geom_text(
    aes(label = percent(prop_high, accuracy = 1)),
    position = position_dodge(width = 0.9),
    vjust = -0.3, size = 4) +
  scale_y_continuous(
    labels = percent_format(accuracy = 1),
    limits = c(0, NA)
  ) +
  scale_fill_manual(
    values = pal_sex
  ) +
  labs(
    title = "Prevalence of high mental health burden ( 14 days) by income and sex",
    x     = "Household income",
    y     = "Proportion with  14 mentally unhealthy days",
    fill  = "Sex"
  ) +
  theme(axis.text.x = element_text(angle = 15, hjust = 1))
```

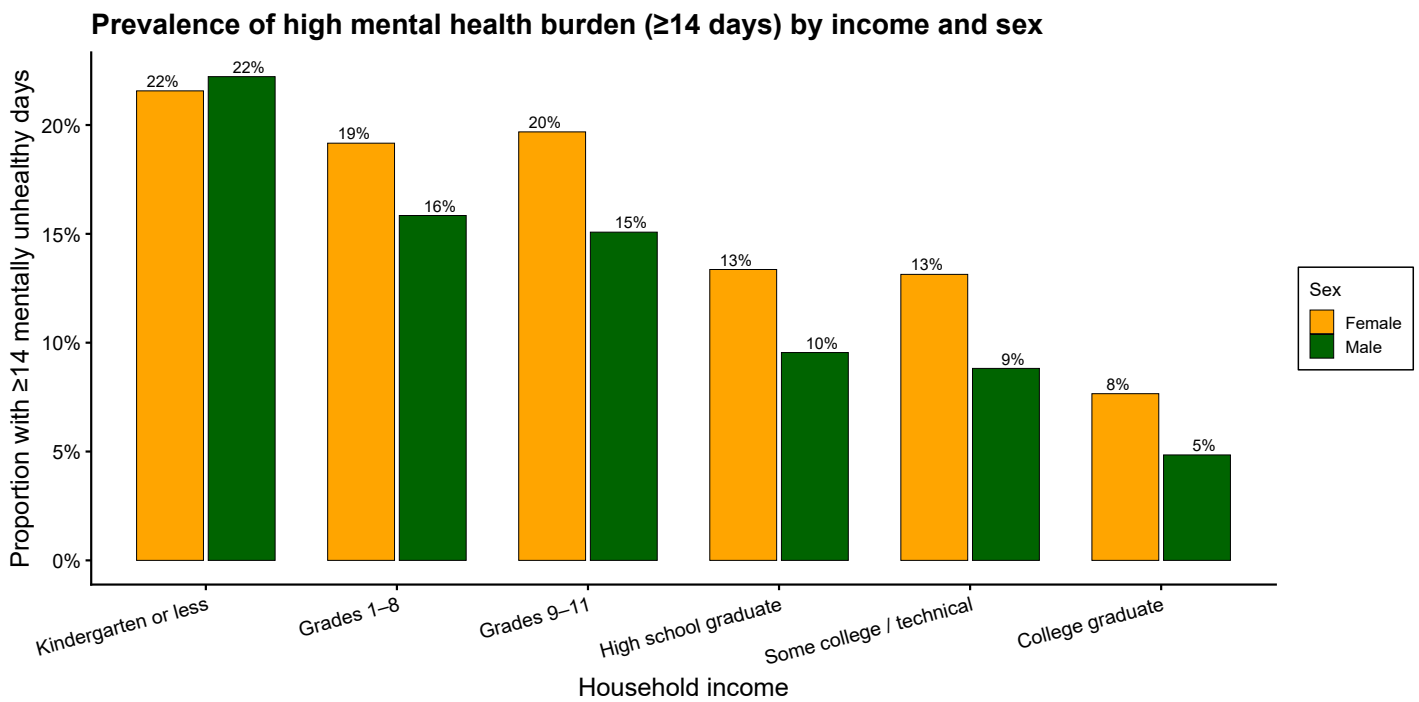

A clear educational gradient is visible in the prevalence of high mental health burden. Individuals with lower levels of educational attainment report substantially higher proportions of ≥14 mentally unhealthy days compared with those with higher education. For example, among respondents with kindergarten or less education, approximately 22% report high mental health burden, whereas this proportion declines to around 8% for females and 5% for males among college graduates.

Across all education categories, females consistently report higher levels of mental health burden than males, although the difference between sexes is relatively modest compared with the overall education gradient. The steady decline in prevalence with increasing educational attainment suggests that socioeconomic resources and opportunities associated with education may play an important role in mental health outcomes.

Overall, the figure indicates that lower educational attainment is associated with substantially higher levels of frequent psychological distress, reinforcing the broader pattern of socioeconomic inequalities in mental health observed in the income analysis.

3.2.4 Mental Health Burden by Age Group and Sex

```
mh_age_sex <- diab |>
  group_by(Age, Sex) |>
  summarise(mean_mh = mean(MentHlth, na.rm = TRUE), .groups = "drop")

pos_dodge <- position_dodge(width = 0.6)

ggplot(mh_age_sex, aes(x = Age, y = mean_mh, colour = Sex)) +
  geom_linerange(
    aes(ymin = 0, ymax = mean_mh),
    position = pos_dodge,
    linewidth = 1
  ) +
  geom_point(
    position = pos_dodge,
    size = 3.8
  ) +
  scale_y_continuous(
    limits = c(0, NA),
```

```
    expand = expansion(mult = c(0, 0.05))
  ) +
  scale_colour_manual(
    values = pal_sex
  ) +
  labs(
    title  = "Mean mentally unhealthy days by age group and sex",
    x      = "Age group",
    y      = "Mean mentally unhealthy days (past 30 days)",
    colour = "Sex"
  )
```

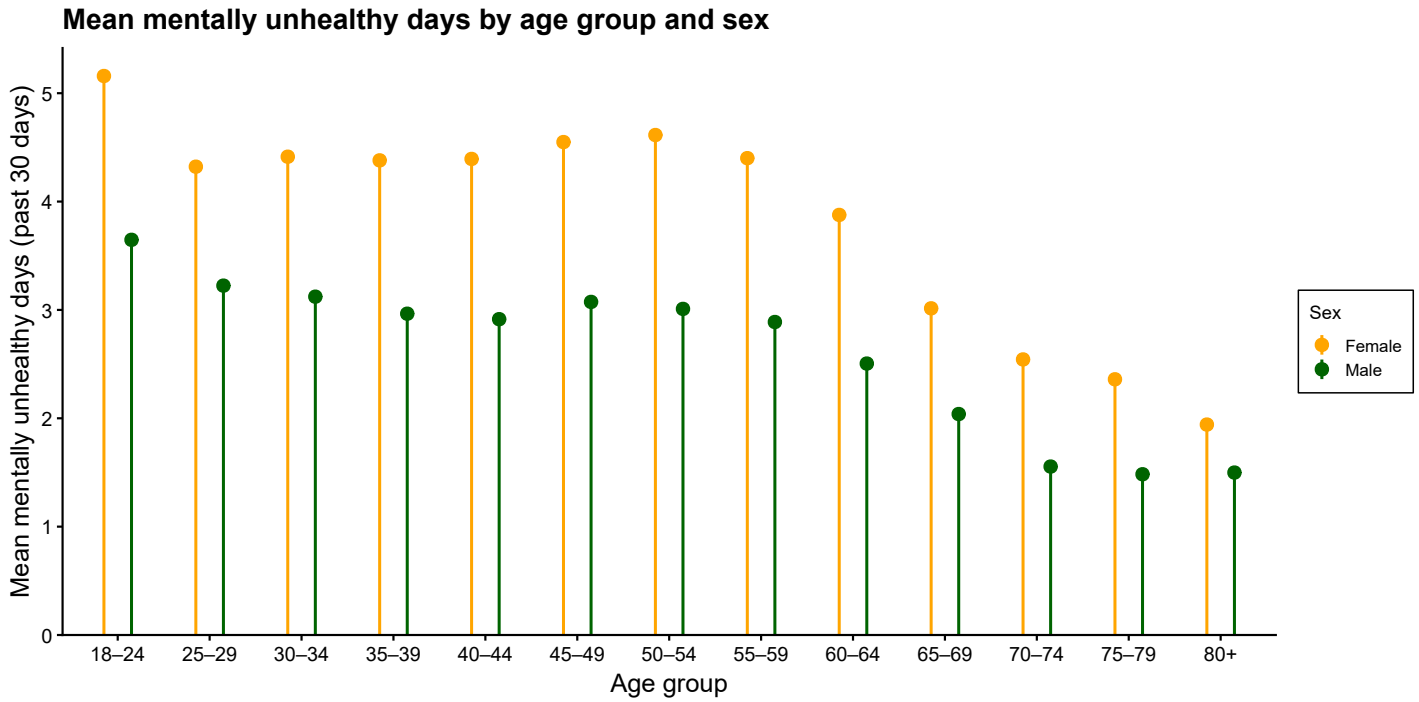

Mean mentally unhealthy days are higher among women than men in every age group. The difference is most visible in younger and middle adulthood, where women report roughly 1–2 more days on average than men.

Across age, the pattern is not monotonic. Mean days are relatively high in early adulthood, remain similar through midlife, and then decline progressively after around age 60 for both sexes. The lowest averages are observed in the oldest age groups.

3.2.5 Joint Distribution of Mental and Physical Health Burden by Sex

```
ggplot(diab, aes(x = MentHlth, y = PhysHlth)) +
  geom_hex(bins = 30) +
  scale_fill_viridis_c(option = "magma", name = "Count", trans = "log10") +
  facet_wrap(~ Sex, nrow = 1) +
  labs(
    title = "Joint distribution of mentally and physically unhealthy days by sex",
    x     = "Mentally unhealthy days (past 30)",
    y     = "Physically unhealthy days (past 30)"
  )
```

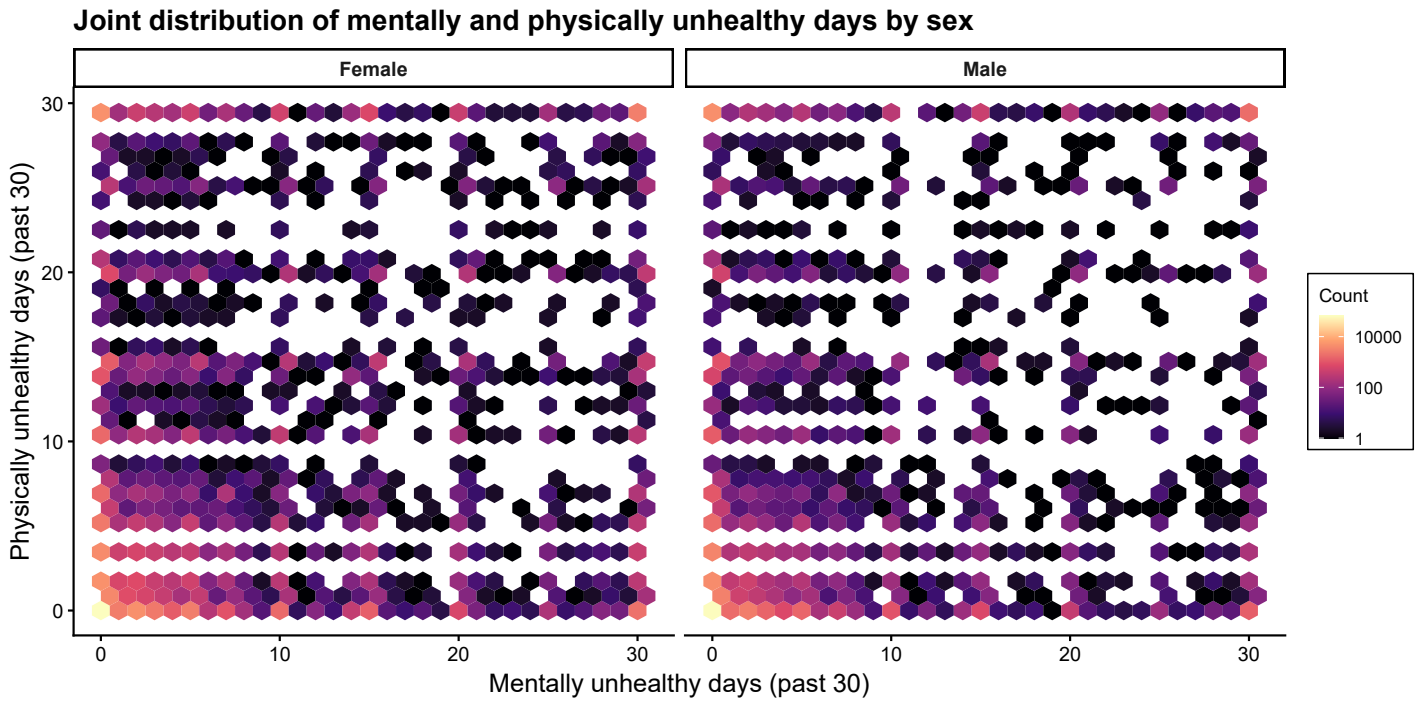

In both panels, the highest concentration of observations occurs near zero mentally and physically unhealthy days. Density decreases as values increase along either axis, and combinations of simultaneously high mental and physical burden are relatively sparse.

There is visible clustering along diagonal regions, indicating that higher numbers of mentally unhealthy days tend to coincide with higher numbers of physically unhealthy days rather than being concentrated along a single axis.

The overall distributional shape is similar for women and men, with broadly comparable density patterns across low, moderate, and higher values.

3.2.6 Prevalence of High Mental Health Burden by Physical and Cardiometabolic indicators

```
diab_mh_phys |>
  ggplot(aes(x = Status, y = prop, fill = Status)) +
  geom_col(width = 0.7, show.legend = FALSE) +
  geom_text(
    aes(label = percent(prop, accuracy = 1)),
    vjust = -0.3,
    size  = 3.3
  ) +
  facet_wrap(~ Indicator, nrow = 2, ncol = 4) +
  scale_y_continuous(
    labels = percent_format(accuracy = 1),
    limits = c(0, NA),
    expand = expansion(mult = c(0, 0.10))
  ) +
  scale_fill_manual(values = pal_binary) +
  labs(
    title = "Prevalence of high Mental Health Burden by Physical and Cardiometabolic Indicators",
    x     = "Indicator status",
    y     = "Proportion with  14 mentally unhealthy days"
  )
```

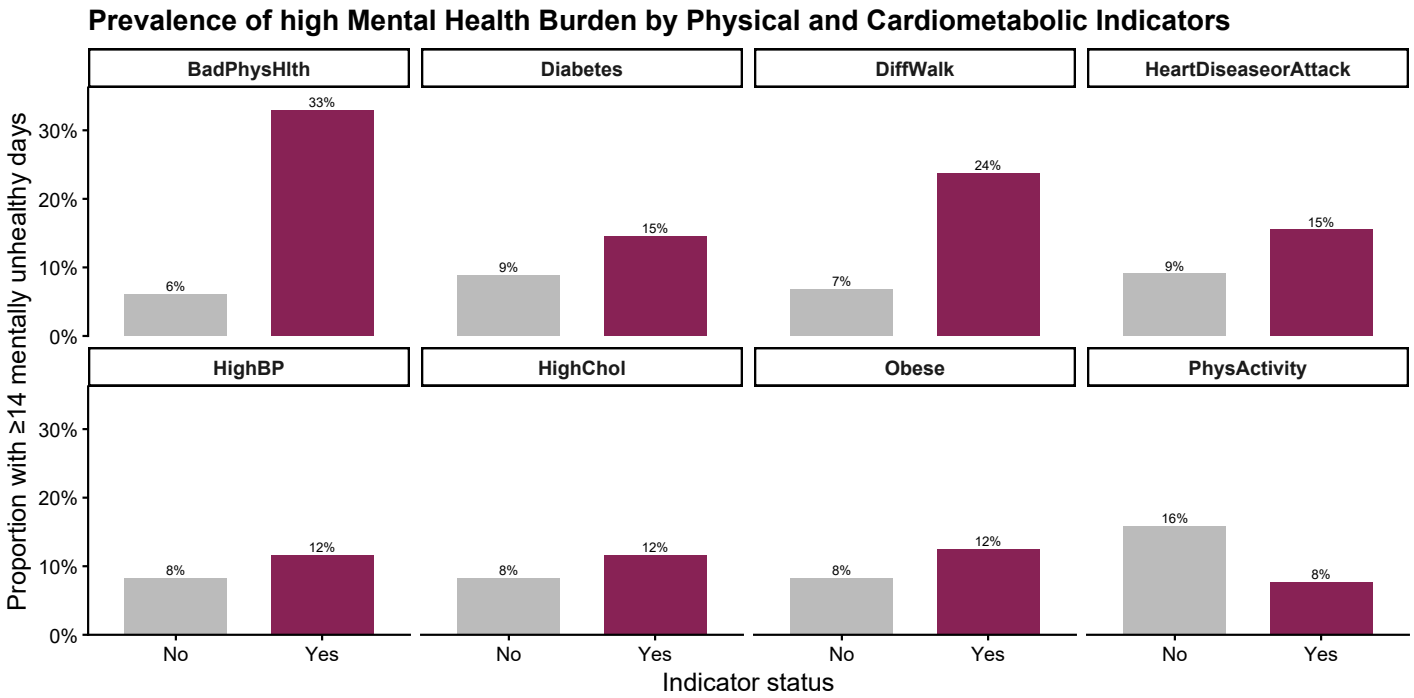

High mental health burden (≥14 mentally unhealthy days) varies substantially across physical and cardiometabolic indicators. The strongest associations are observed for indicators of functional and physical health limitations. Individuals reporting high physical health burden (≥14 physically unhealthy days) show a markedly elevated prevalence of high mental health burden (≈ 33%), compared with only about 6% among those without high physical burden. A similarly large difference appears for difficulty walking, where prevalence rises from roughly 7% to 24%.

Cardiometabolic conditions such as diabetes, heart disease or myocardial infarction, high blood pressure, and high cholesterol are also associated with higher levels of mental health burden, although the magnitude of these differences is smaller. For example, respondents with diabetes or heart disease exhibit prevalence around 15%, compared with approximately 9% among those without these conditions.

Lifestyle factors show a mixed pattern. Obesity is associated with moderately higher mental health burden (≈12% vs ≈8%), while physical activity appears protective, with prevalence roughly halved among individuals reporting recent physical activity (≈ 8%) compared with those reporting none (≈16%).

Overall, the results suggest that mental health burden is most strongly linked with physical health limitations and functional impairment, while cardiometabolic disease and lifestyle factors exhibit weaker but still noticeable associations.

3.2.7 Summary of Key Findings

Several clear structural patterns emerge in the analysis of mental health burden. First, the distribution of mentally unhealthy days is highly right-skewed, with a large concentration at zero and a long tail toward higher values. This indicates substantial heterogeneity in psychological distress across the population and motivates the use of count-based models with potential zero inflation in subsequent modelling.

Second, strong **socioeconomic gradients** are visible. The prevalence of high mental health burden declines consistently with increasing income and education, suggesting that socioeconomic disadvantage is associated with substantially higher levels of frequent psychological distress.

Third, **demographic differences** are evident. Women report higher average numbers of mentally unhealthy days than men across all age groups, while mental health burden tends to decline gradually in older age groups after midlife.

Fourth, mental and physical health burdens show **moderate co-occurrence**, with higher levels of mentally unhealthy days tending to coincide with higher levels of physically unhealthy days.

Finally, mental health burden is **most strongly associated with indicators of physical limitation**, particularly poor physical health and difficulty walking. Cardiometabolic conditions such as diabetes, heart disease, hypertension, and high cholesterol also show positive associations with mental health burden, although these relationships are weaker than those linked to functional impairment.

Overall, the results suggest that mental health burden in this population is structured primarily by **socioeconomic position, physical health status, and functional limitation**, providing a clear empirical basis for the modelling analyses developed in later sections.

3.2.8 Implications for Modelling

The exploratory analysis provides several insights relevant for model development. The distribution of mentally unhealthy days is highly right-skewed with a substantial concentration at zero, suggesting that standard linear modelling would be inappropriate and motivating the use of count-based models such as Poisson or negative binomial regression, potentially with zero-inflation.

Socioeconomic gradients observed across income and education indicate that demographic and socioeconomic variables should be included as key predictors.

Finally, the strong co-occurrence between mental and physical health indicators suggests that physical health variables such as poor physical health days, difficulty walking, and chronic disease indicators may play an important explanatory role in modelling mental health burden.

3.x.x Spearman Correlation Heatmap: Continuous and Ordinal Variables

```r
ggplot(spear_tidy, aes(x = Var1, y = Var2, fill = rho)) +
  geom_tile(colour = "white") +
  geom_text(
    aes(label = sprintf("%.2f", rho)),
    size   = 3.5,
    colour = "black"
  ) +
  scale_fill_gradient2(
    low      = "#2166AC",
    mid      = "white",
    high     = "#D6604D",
    midpoint = 0,
    limits   = c(-1, 1),
    name     = "Spearman's \u03c1"
  ) +
  labs(
    title = "Spearman rank correlations among continuous and ordinal variables",
    x     = NULL,
    y     = NULL
  ) +
  theme(
    axis.text.x = element_text(angle = 40, hjust = 1),
    axis.text.y = element_text(hjust = 1)
  )
```

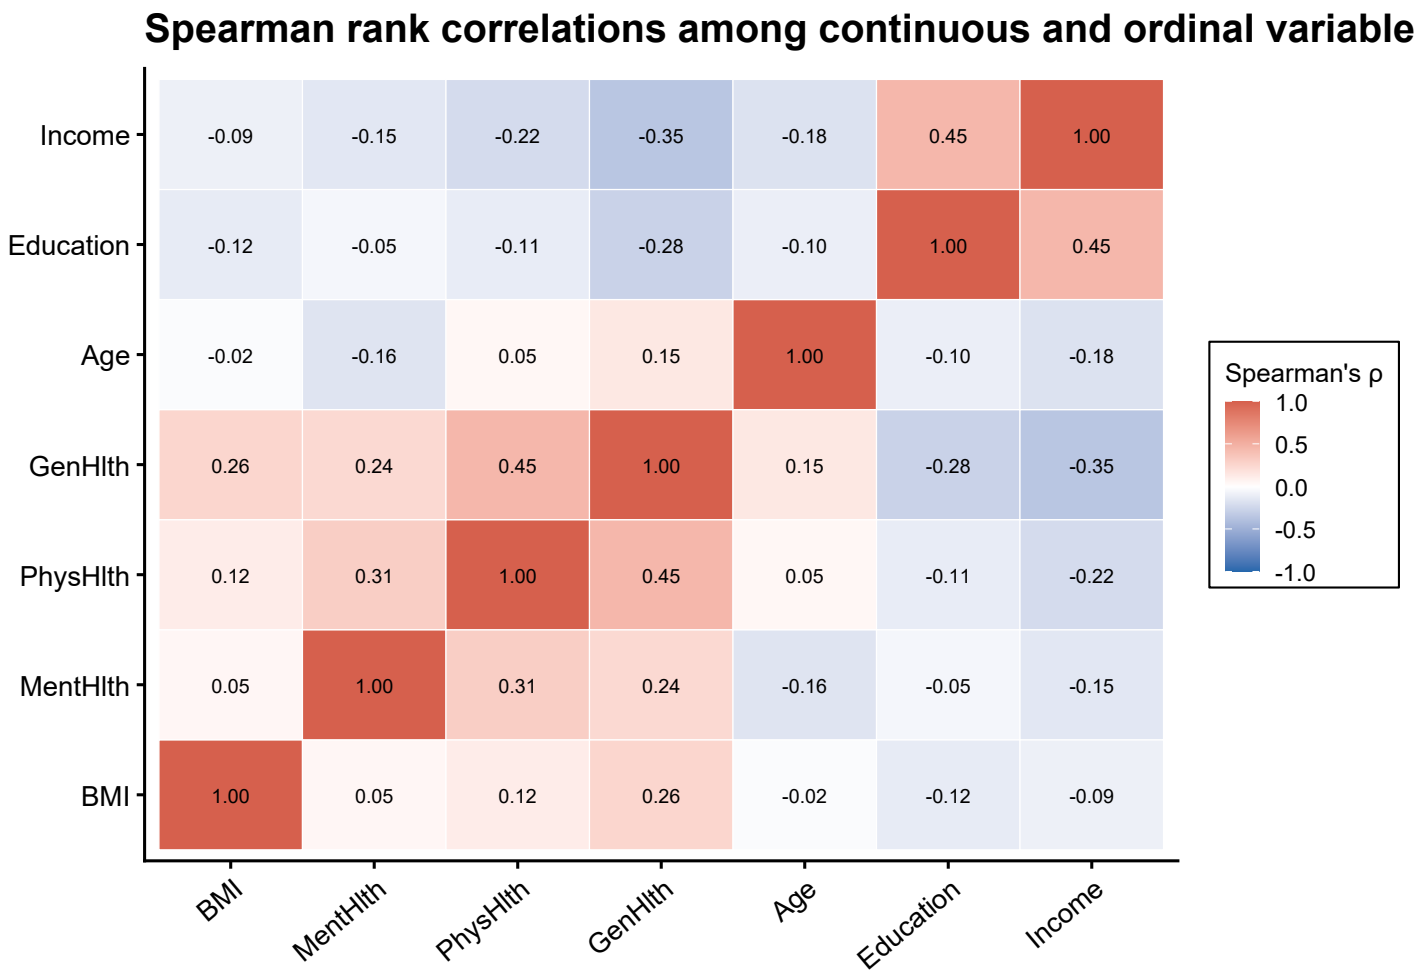

The Spearman rank correlations indicate mostly **weak to moderate** monotonic associations among the continuous and ordinal variables.

- The strongest positive correlations are between **GenHlth and PhysHlth** ($\rho \approx 0.45$) and between **Education and Income** ($\rho \approx 0.45$), suggesting moderate monotonic relationships in both cases.
- **GenHlth** is moderately **negatively** correlated with **Income** ($\rho \approx -0.35$) and **Education** ($\rho \approx -0.28$). Because higher GenHlth scores denote poorer health, these negative correlations imply that higher income and education are associated with **better** self-rated health.
- **MentHlth** and **PhysHlth** show a modest positive association ($\rho \approx 0.31$), consistent with some co-occurrence of mental and physical health burden.
- Most other pairwise correlations are small ($|\rho| < 0.3$). In particular, **BMI** and **Age** have only weak associations with the other variables in this set.

Overall, the correlation matrix does **not** indicate problematic collinearity: apart from a few moderate associations (GenHlth–PhysHlth, Education–Income), most variables are only weakly related. This supports their concurrent use as predictors in multivariable models, subject to routine multicollinearity checks.

3.x.x Cramér's V Association Heatmap: Binary Clinical and Lifestyle Variables

To complement the Spearman correlation heatmap for continuous/ordinal variables, the strength of pairwise associations among all binary clinical and lifestyle indicators is quantified using Cramér's V. This symmetric measure summarises the association in each 2×2 contingency table on a 0–1 scale, without assuming any directionality. Diabetes status is deliberately excluded.

Together, Sections 3.14 and 3.15 establish the correlation structure among all key variables, identifying which predictors are inter-related and justifying the need for multivariable adjustment in all subsequent models.

```r
ggplot(cv_tidy, aes(x = Var1, y = Var2, fill = V)) +
  geom_tile(colour = "white") +
  geom_text(
    aes(label = sprintf("%.2f", V)),
    size   = 3.8,
    colour = "black"
  ) +
  scale_fill_gradient2(
    low       = "white",
    high      = "#D6604D",
    midpoint = 0,
    limits   = c(0, 1),
    name     = "Cramér's V"
  ) +
  labs(
    title = "Cramér's V among binary clinical and lifestyle variables",
    x = NULL,
    y = NULL) +
  theme(axis.text.x = element_text(angle = 30, hjust = 1))
```

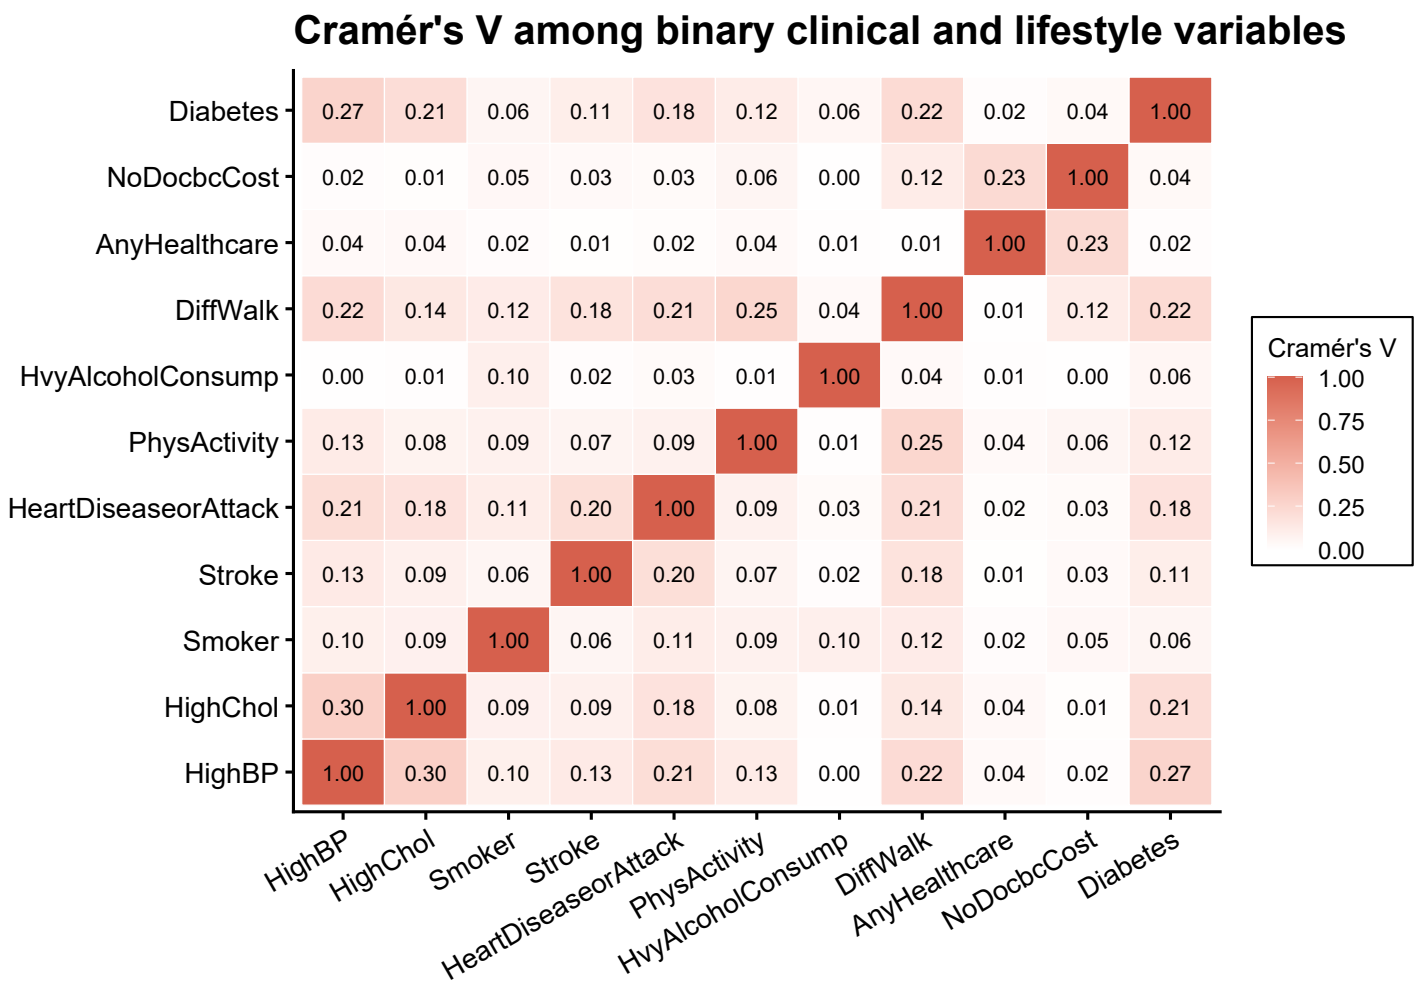

The Cramér's V matrix indicates that most pairwise associations among the binary clinical and lifestyle variables are **weak**, with the majority of values below 0.20.

The strongest association is between **HighBP and HighChol** (V ≈ 0.30), representing a moderate relationship relative to the rest of the matrix. This aligns with known clustering of cardiovascular risk factors.

Other comparatively larger associations include:

- **PhysActivity and DiffWalk** (V ≈ 0.25)
- **HighBP and DiffWalk** (V ≈ 0.22)
- **HighBP and HeartDiseaseorAttack** (V ≈ 0.21)
- **Stroke and HeartDiseaseorAttack** (V ≈ 0.20)

These moderate associations are clinically plausible and reflect shared cardiovascular or functional impairment pathways.

Lifestyle indicators such as **Smoker** and **HvyAlcoholConsump** exhibit generally weak associations with other variables (typically V < 0.15). Healthcare access indicators (**AnyHealthcare**, **NoDocbcCost**) also show minimal association with most clinical variables.

Analytical Implication

The matrix does not suggest substantial redundancy among the binary predictors. Aside from a small number of moderate associations (primarily within cardiovascular-related variables), most relationships are weak.

This supports inclusion of these variables simultaneously in multivariable models, with routine collinearity diagnostics (e.g., VIF or condition indices) as a precaution rather than an expected issue.

4. Refined Research Questions

RQ1: Multidomain Determinants and Prediction of Diabetes Status

Outcome Variable: *Diabetes — Multiclass Outcome*

Research Question

> *To what extent can diabetes status be explained and predicted using demographic, socioeconomic, behavioural, and cardiometabolic factors, and which predictors contribute most strongly to distinguishing between no diabetes, prediabetes, and diabetes?*

Rationale

Exploratory analysis of the CDC Diabetes Health Indicators dataset indicates that diabetes status is structured by multiple interacting determinants. Section 3.2.1 demonstrated a strong **age gradient**, with prediabetes and diabetes concentrated in older age groups, suggesting that demographic factors must be accounted for when evaluating other risk indicators.

Cardiometabolic clustering was evident across several visualisations. The joint distribution of **high blood pressure and high cholesterol** increased substantially across diabetes categories, while the risk-factor accumulation plot showed a clear gradient in which diabetes prevalence rises with the number of cardiometabolic risk indicators. This pattern aligns with established evidence linking obesity, hypertension, and dyslipidaemia to type 2 diabetes risk (Ismail et al., 2021).

Behavioural factors also exhibited variation across diabetes categories. Physical activity levels tended to decline with worsening diabetes status, while smoking and alcohol consumption showed weaker separation. Because behavioural variables are potentially modifiable, evaluating their contribution to diabetes risk is relevant for prevention strategies (Aune et al., 2015).

Finally, socioeconomic gradients were visible in the education–income distribution, where diabetes prevalence was higher among lower socioeconomic groups. These patterns suggest that diabetes risk is embedded within broader social determinants of health, which influence exposure to risk factors and access to care (Hill-Briggs et al., 2021).

Together, these findings motivate a multivariable modelling approach to estimate adjusted associations and evaluate predictive performance across diabetes categories.

'EDI Discussed in Task1'

RQ2: Predicting Cardiovascular Events Using Machine Learning

Outcome Variables: *HeartDiseaseorAttack* (primary), *Stroke* (secondary) — binary outcomes

Research Question

> *How accurately can cardiovascular events (heart disease or heart attack, and separately stroke) be predicted using demographic, behavioural, and clinical risk factors, and which variables contribute most strongly to individual-level risk predictions as explained by SHAP values?*

Rationale

Exploratory analysis identified strong associations between several risk factors and cardiovascular outcomes. Age showed the most pronounced gradient, with the prevalence of both heart disease and stroke increasing substantially across older age groups. Additional bivariate analyses indicated consistent relationships between cardiovascular events and established clinical and behavioural indicators, including **high blood pressure, high cholesterol, body mass index, smoking status, and physical inactivity**.

These predictors are known to interact in complex and potentially nonlinear ways. Traditional regression models require explicit specification of such interactions, whereas **Random Forest** models can capture nonlinear relationships and high-order interactions automatically through ensemble decision trees. This makes the method well suited for analysing large observational datasets with mixed predictor types and complex dependency structures.

Although Random Forest models often achieve strong predictive performance, they are typically difficult to interpret. To address this limitation, **SHapley Additive exPlanations (SHAP)** are used to quantify the contribution of each predictor to individual model predictions. SHAP values provide both **global feature importance** and **local explanations**, allowing identification of which factors increase or decrease predicted cardiovascular risk for specific individuals (Lundberg & Lee, 2017).

The modelling approach therefore combines predictive accuracy with interpretability, enabling assessment of both **which variables are most influential in predicting cardiovascular events** and **how these variables influence individual-level predictions**.

EDI / Ethical Considerations

Machine learning models used in health prediction must be interpretable and transparent. Using SHAP allows examination of whether predictions rely disproportionately on demographic characteristics such as age or sex, helping identify potential biases. Evaluating model performance across socioeconomic groups is also important to ensure that predictive accuracy does not vary systematically across population subgroups, which could otherwise reinforce existing health inequalities.

RQ3: Socioeconomic and Lifestyle Determinants of Mental Health Burden

Outcome Variable: *MentHlth* — number of self-reported mentally unhealthy days in the past 30 days (0–30)

Research Question

> *To what extent do income, education, physical activity, smoking, and heavy alcohol consumption predict the number of mentally unhealthy days reported in the past 30 days, and does the association between physical activity and mental health burden differ across income groups?*

Rationale

The *MentHlth* variable represents a count of mentally unhealthy days within the past month (0–30). Exploratory analysis revealed a strongly right-skewed distribution with a large concentration of zero values, indicating that many respondents report no mentally unhealthy days while a smaller subset experience substantial burden. Such distributions violate assumptions of linear models and often produce overdispersion and excess zeros, making standard Poisson or logistic approaches unsuitable.

A Zero-Inflated Negative Binomial (ZINB) model is therefore appropriate because it separates two processes:

1. the probability that individuals experience no mental health burden, and
2. the frequency of mentally unhealthy days among those who do report burden (Rust et al., 2014; Liu et al., 2023).

The EDA further revealed clear socioeconomic gradients in mental health outcomes. Respondents in the lowest income categories exhibited substantially higher prevalence of high mental health burden compared with those in higher income groups. Physical activity also showed a consistent inverse relationship with mentally unhealthy days, suggesting potential behavioural protection.

These patterns motivate a modelling approach that jointly evaluates socioeconomic position and lifestyle behaviours. In particular, testing an interaction between income and physical activity allows examination of whether the protective association of physical activity differs across socioeconomic strata.

Prior research supports the relevance of these predictors. Socioeconomic disadvantage has been consistently linked with elevated risk of mental health disorders and poorer psychological well-being (Kessler et al., 2010), while analyses of BRFSS data have shown that ZINB regression provides an effective framework for modelling unhealthy-days indicators with excess zeros (Rust et al., 2014; Liu et al., 2023).

EDI / Ethical Considerations

Socioeconomic gradients in mental health outcomes raise important equity considerations. If income remains strongly associated with mental health burden after adjustment for behavioural factors, this suggests that structural determinants, such as economic insecurity and access to resources, play a major role in shaping mental health outcomes. Understanding these relationships can inform more equitable public health strategies that address structural conditions rather than focusing solely on individual behaviour.

6. Multidomain Determinants and Prediction of Diabetes Status

6.1 Modelling Objective

This section addresses the following research question:

> To what extent can diabetes status be explained and predicted using demographic, socioeconomic, behavioural, and cardiometabolic factors, and which predictors contribute most strongly to distinguishing between no diabetes, prediabetes, and diabetes?

The objective is twofold:

1. **Explanatory** – to estimate the adjusted associations between predictor domains and diabetes status.
2. **Predictive** – to compare the classification performance of statistical and machine learning models.

Diabetes status is treated as a **multiclass outcome** (No diabetes / Prediabetes / Diabetes), requiring modelling approaches suitable for nominal categorical outcomes.

6.2. Outcome Variable

The dependent variable is **diabetes status**, categorised into:

- **No diabetes**
- **Prediabetes**
- **Diabetes.**

Exploratory analysis demonstrated class imbalance, with the majority of respondents in the no-diabetes category and substantially smaller proportions in the prediabetes and diabetes groups. This imbalance is accounted for in model training and evaluation.

6.3 Training and Validation

6.3.1 Data Partitioning and Resampling:

The dataset is partitioned into a 75% training set and a 25% testing set using stratified sampling on diabetes status. Within the training set, 5-fold stratified cross-validation is implemented to estimate model performance during development. The held-out test set is reserved for final evaluation to ensure unbiased performance assessment.

```r
set.seed(42)

diabetes_split <- initial_split(
  diab,
  prop = 0.75,
  strata = Diabetes
)

diab_train <- training(diabetes_split)
diab_test  <- testing(diabetes_split)


diabetes_folds <- vfold_cv(
  diab_train,
  v = 5,
  strata = Diabetes
)
```

6.3.2 Recipe Definition and Preprocessing Strategy

- **Outcome specification** – Models `Diabetes` as a three-level categorical outcome (No diabetes, Prediabetes, Diabetes).
- **Near-zero variance removal** – Removes predictors with minimal variability to prevent instability and redundant information.
- **Ordinal scoring** – Converts ordered factors (`Age`, `Education`, `Income`, `GenHlth`) into numeric scores while preserving their natural ranking. This introduces an assumption of approximately equal spacing between adjacent categories; however, the primary objective is to retain monotonic structure in a computationally efficient representation suitable for both parametric and tree-based models.
- **Dummy encoding** – Transforms remaining nominal predictors into binary indicator variables suitable for modelling.
- **Normalization** – Standardizes numeric predictors (mean = 0, sd = 1) to improve optimisation and comparability across models.
- **SMOTE balancing** – Applies synthetic oversampling to minority classes within resampling folds to mitigate severe class imbalance.

```r
rec_rq1 <- recipe(Diabetes ~ ., data = diab_train) |>
  step_nzv(all_predictors()) |>
  step_ordinalscore(Age, Education, Income, GenHlth) |>
  step_dummy(all_nominal_predictors()) |>
  step_normalize(BMI, MentHlth, PhysHlth) |>
  step_smote(Diabetes, over_ratio = 0.5, seed = 42)

rec_rq1
```

6.3.3 Model Specification

We now define the candidate models for RQ1 using the **tidymodels** framework.

Each model is specified independently of preprocessing and will later be combined with the common recipe using **workflows** and **workflow sets**.

Model	Type	Primary Function	Key Strengths
Multinomial Logistic Regression	Parametric	Linear Baseline	Interpretable log-odds; direct coefficient analysis.

Random Forest	Ensemble (Bagging)	Non-linear Discovery	Captures complex interactions; robust to multicollinearity.
XGBoost	Ensemble (Boosting)	Predictive Performance	High efficiency on tabular data; optimized for accuracy.

```
# Multinomial logistic regression
mod_multinom <-
  multinom_reg() |>
  set_engine("nnet") |>
  set_mode("classification")

# Random Forest
mod_rf <-
  rand_forest(
    trees = 256,
    mtry  = tune(),
    min_n = tune()
  ) |>
  set_engine("ranger") |>
  set_mode("classification")

# Gradient Boosted Trees
mod_xgb <-
  boost_tree(
    trees = 512,
    tree_depth = tune(),
    learn_rate = tune(),
    loss_reduction = tune(),
    mtry = tune(),
    min_n = tune()
  ) |>
  set_engine("xgboost") |>
  set_mode("classification")
```

6.3.4 Workflow Construction and Workflow Set Definition

We now integrate:

- The **RQ1 preprocessing recipe**
- The **candidate model specifications**

  using `workflow()` and `workflow_set()`.

Each workflow represents one **model–preprocessing pair**.
The workflow set enables parallel evaluation of all candidate models under identical resampling conditions.

```
# Individual workflows
wf_multinom <-
  workflow() |>
  add_recipe(rec_rq1) |>
  add_model(mod_multinom)

wf_rf <-
  workflow() |>
```

```
  add_recipe(rec_rq1) |>
  add_model(mod_rf)

wf_xgb <-
  workflow() |>
  add_recipe(rec_rq1) |>
  add_model(mod_xgb)
```

```
# Workflow set
wf_set_rq1 <-
  workflow_set(
    preproc = list(rq1_recipe = rec_rq1),
    models  = list(
      multinom = mod_multinom,
      rf       = mod_rf,
      xgb      = mod_xgb))
```

6.3.5 Hyperparameter Tuning and Resampled Model Fitting

Metrics computed

Model performance is evaluated using multiple complementary metrics to account for class imbalance and provide a robust assessment of predictive performance:

- **Accuracy** – overall classification rate
- **Balanced Accuracy** – mean recall across classes, correcting for class imbalance
- **Cohen's Kappa** – agreement between predictions and observed classes beyond chance
- **F1-score (macro)** – harmonic mean of precision and recall averaged across classes
- **ROC-AUC (multiclass)** – ability of the model to discriminate between diabetes categories

```
metrics_diab <-
  metric_set(
    accuracy,
    bal_accuracy,
    kap,
    f_meas,
    roc_auc
  )
```

We now estimate model performance using the 5-fold cross-validation object (`diab_folds`).

- Logistic regression → no tuning (baseline model)
- Random Forest → tune `mtry`, `min_n`
- XGBoost → tune depth, learning rate, regularisation, etc.

Hyperparameter tuning was conducted using grid search across 5 candidate configurations for each model. This grid size provides a balance between exploration of the hyperparameter space and computational feasibility given the large dataset.

```
# set.seed(42)
# ctrl <- control_grid(
#   save_pred = TRUE,
#   save_workflow = FALSE,
#   verbose = TRUE
# )
#
# wf_set_rq1_fit <-
```

```r
#   wf_set_rq1 |>
#   workflow_map(
#     "tune_grid",
#     resamples = diabetes_folds,
#     grid = 10,
#     metrics = metrics_diab,
#     control = ctrl
#     )

wf_set_rq1_fit <- readRDS('wf_set_rq1_fit.rds')
```

```r
m1 <- pretty_autoplot(
  wf_set_rq1_fit,
  "roc_auc",
  "ROC-AUC (multiclass)",
  "Discrimination (ROC-AUC)")
m2 <- pretty_autoplot(
  wf_set_rq1_fit,
  "bal_accuracy",
  "Balanced accuracy",
  "Class-balanced accuracy")
m3 <- pretty_autoplot(
  wf_set_rq1_fit,
  "f_meas",
  "Macro F1-score",
  "Macro F1 (imbalance-aware)")
m4 <- pretty_autoplot(wf_set_rq1_fit, "kap", "Cohen's kappa", "Agreement beyond chance")

(m1 | m2) / (m3 | m4) +
  plot_layout(guides = "collect") &
  theme(legend.position = "bottom")
```

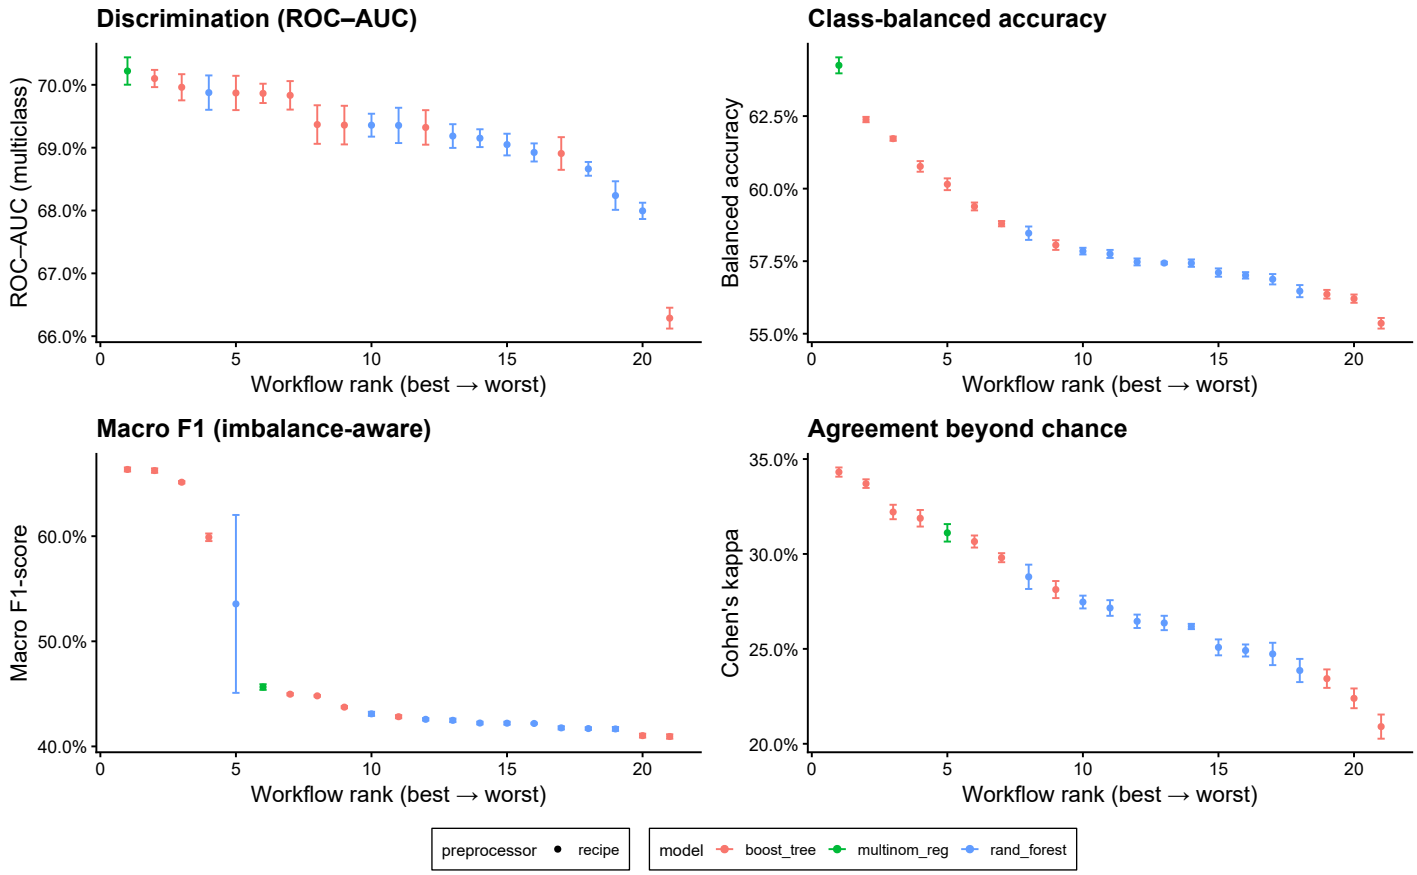

ROC–AUC.
ROC–AUC measures a model's ability to discriminate between the three diabetes categories across all possible classification thresholds. The highest ROC–AUC value is achieved by the **multinomial logistic regression model**, which appears as the top-ranked workflow in this metric. However, beyond this leading configuration, most of the remaining high-ranking workflows consist of **XGBoost and random forest models**, with values clustered between approximately 0.69 and 0.70. This indicates that while logistic regression achieves the single best ROC–AUC score, tree-based models provide similarly competitive discrimination across multiple configurations.

Balanced Accuracy.
Balanced accuracy reflects the average recall across classes, giving equal importance to each diabetes category. The **multinomial logistic regression model again appears as the top-ranked workflow**, achieving the highest balanced accuracy among the candidate models. Several **XGBoost configurations follow closely**, while random forest models occupy the middle range of the ranking. Although performance differences are relatively modest, the results suggest that multiple model types are capable of achieving comparable class-balanced recall.

Macro F1 Score.
Macro F1 was selected as the primary tuning metric because it balances precision and recall equally across classes and is therefore appropriate under class imbalance. In this panel, the **top-ranked workflows are clearly dominated by XGBoost models**, which achieve substantially higher macro F1 scores than the other model types. Random forest configurations perform considerably lower, and the multinomial logistic regression model appears further down the ranking. This pattern indicates that boosted trees are better able to balance precision and recall across the three diabetes categories.

Cohen's Kappa.
Cohen's kappa measures agreement between predicted and observed classifications beyond chance. The highest kappa values are achieved by **XGBoost configurations**, which occupy several of the top-ranked positions in this panel. The multinomial logistic regression model appears among the higher-performing workflows but does not exceed the best boosted-tree configurations. Random forest models generally show lower kappa values and appear further down the ranking. Overall, this metric again indicates stronger agreement for boosted tree models compared with the other modelling approaches.

6.3.6 Extract tuning results

```
# Extract tuning results
rf_res  <- extract_workflow_set_result(wf_set_rq1_fit, "rq1_recipe_rf")
xgb_res <- extract_workflow_set_result(wf_set_rq1_fit, "rq1_recipe_xgb")

# Select best hyperparameters using Macro F1
best_rf  <- select_best(rf_res, metric = "f_meas")
best_xgb <- select_best(xgb_res, metric = "f_meas")

# Extract workflows
rf_wf       <- extract_workflow(wf_set_rq1_fit, "rq1_recipe_rf")
xgb_wf      <- extract_workflow(wf_set_rq1_fit, "rq1_recipe_xgb")
multinom_wf <- extract_workflow(wf_set_rq1_fit, "rq1_recipe_multinom")

# Finalise tuned workflows
rf_final_wf  <- finalize_workflow(rf_wf, best_rf)
xgb_final_wf <- finalize_workflow(xgb_wf, best_xgb)

# Multinomial requires no tuning, so it is already final
multinom_final_wf <- multinom_wf
```

```
rf_final_fit       <- fit(rf_final_wf, data = diab_train)
xgb_final_fit      <- fit(xgb_final_wf, data = diab_train)
multinom_final_fit <- fit(multinom_final_wf, data = diab_train)
```

6.4 Model Evaluation

6.4.1 ROC Curves

```
autoplot(roc_all) +
  labs(
    title = "Multiclass ROC curves (one-vs-rest) - Test set",
    subtitle = "All three models overlaid; each panel shows one class vs the rest",
    x = "1 - Specificity",
    y = "Sensitivity"
  ) +
  theme(
    legend.position = 'bottom'
  )
```

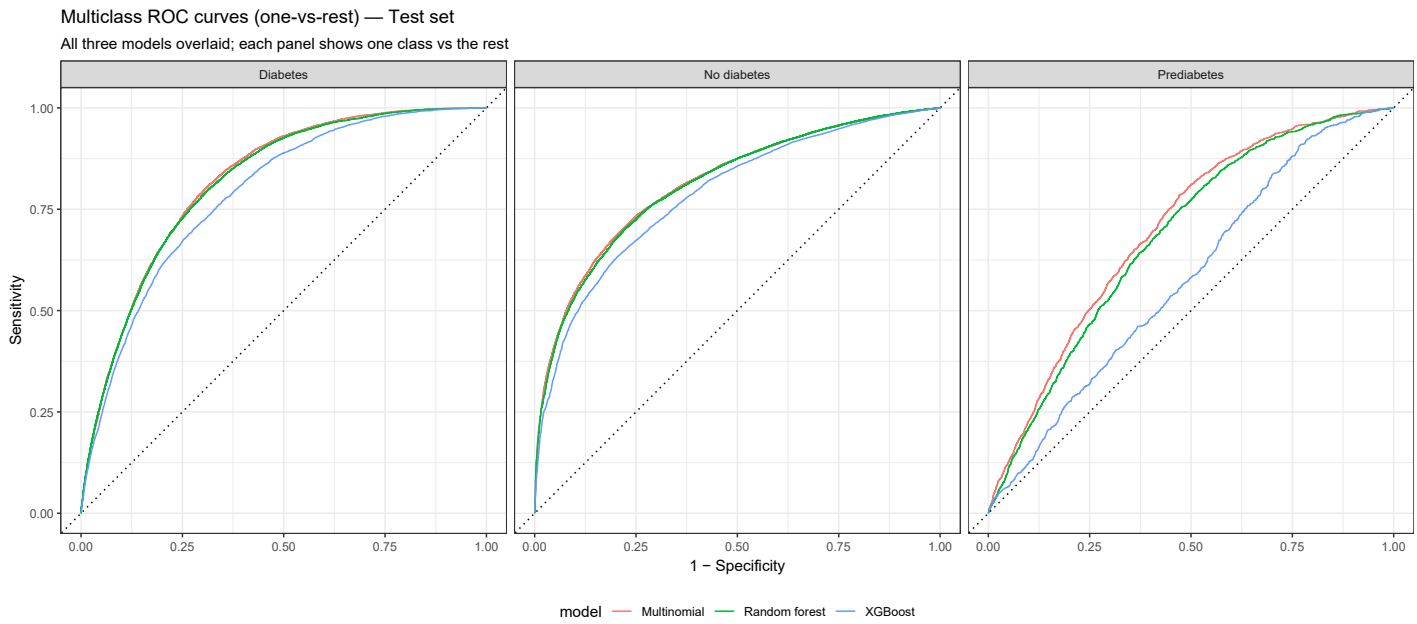

The multiclass ROC curves provide additional insight into discrimination performance across the three classes.

For **diabetes** and **no-diabetes**, all models show strong discriminative ability, with curves well above the diagonal reference line. Differences between models are relatively small, though the tree-based models appear slightly stronger in some regions of the curve.

For the **prediabetes class**, however, the multinomial logistic regression model clearly outperforms both random forest and XGBoost. The ROC curve for the multinomial model remains consistently above the other two models, indicating better sensitivity–specificity trade-offs for detecting this intermediate category.

6.4.2 Confusion Matrices

```
p_cm_xgb <- plot_cm(pred_xgb, "Confusion matrix - XGBoost")
p_cm_rf  <- plot_cm(pred_rf,  "Confusion matrix - Random forest")
p_cm_mn  <- plot_cm(pred_mn,  "Confusion matrix - Multinomial Logistic Regression")

(p_cm_xgb | p_cm_rf) / p_cm_mn +
  plot_layout(guides = "collect")
```

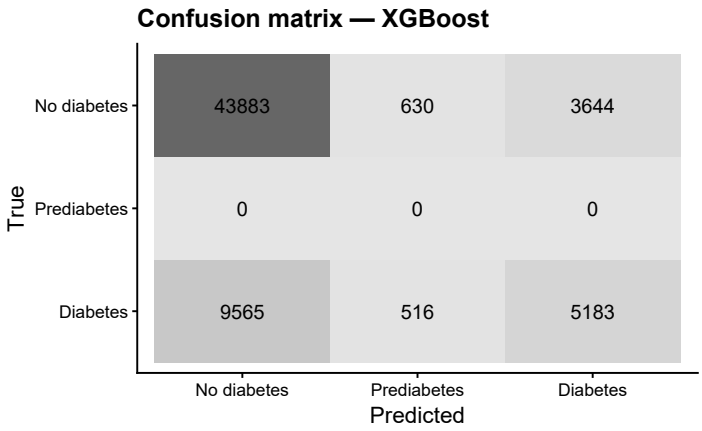

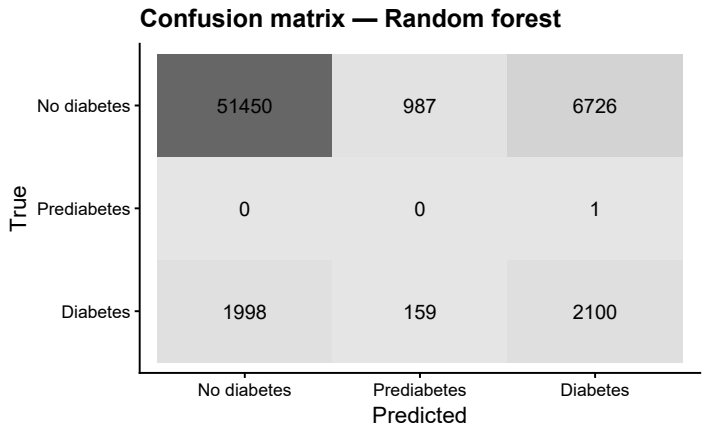

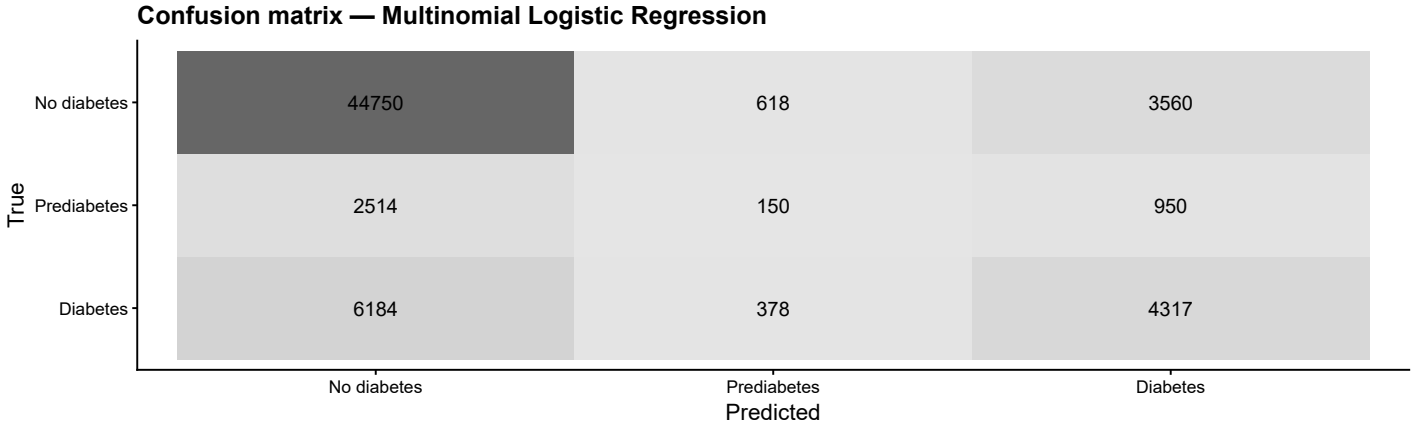

The confusion matrices reveal notable differences in classification behaviour across the three models.

The **multinomial logistic regression model** is the only model that correctly identifies a meaningful number of **prediabetes cases** (150 observations). In contrast, both **XGBoost and random forest completely fail to predict the prediabetes class**, assigning all observations in this category to either the no-diabetes or diabetes classes. This indicates that the tree-based models effectively collapse the problem into a two-class classification task, ignoring the intermediate category.

For the **no-diabetes class**, all three models achieve high numbers of correct predictions due to the dominance of this category in the dataset. However, the multinomial model shows slightly higher correct classifications than the boosted tree model.

For the **diabetes class**, all models correctly identify several thousand cases, though misclassification between diabetes and no-diabetes remains present across all models.

Overall, the multinomial model demonstrates a **more balanced classification behaviour**, particularly by retaining predictive capacity for the minority prediabetes class.

Overall Assessment

Although gradient boosting and random forest performed competitively during cross-validation, the held-out test set evaluation suggests that **multinomial logistic regression** provides the most balanced performance across classes. In particular, it is the only model capable of correctly identifying prediabetes cases while maintaining strong discrimination for the other two classes.

6.5 Model Interpretation

To complement predictive evaluation, model interpretability techniques were used to understand the predictors most strongly associated with diabetes status. Two complementary approaches were used. First, **multinomial logistic regression coefficients were examined to estimate log-odds ratios and odds ratios**, providing interpretable measures of association between predictors and diabetes outcomes. Second, **SHAP (SHapley Additive exPlanations) values were computed for the gradient boosted tree model** to identify the predictors contributing most strongly to model predictions in a non-linear modelling framework.

Together, these approaches provide both **statistical interpretability and machine-learning feature attribution**, enabling a comprehensive understanding of the determinants of diabetes status.

6.5.1 Log-Odds Ratios (Multinomial Logistic Regression)

Multinomial logistic regression estimates the log-odds of belonging to each diabetes category relative to a reference category. In this analysis, **No diabetes** serves as the reference class.

Positive coefficients indicate **higher odds of belonging to the specified diabetes category**, while negative coefficients indicate **lower odds relative to the reference group**.

```
multinom_coefs_tidy |>
  filter(term != "(Intercept)") |>
  ggplot(aes(x = odds_ratio, y = reorder(term, odds_ratio))) +
  geom_point()  +
  geom_vline(xintercept = 1, linetype = "dashed") +
  facet_wrap(~ y.level) +
  labs(
    x = "Odds Ratio",
    y = "Predictor",
    title = "Odds ratios for predictors of diabetes status"
  )
```

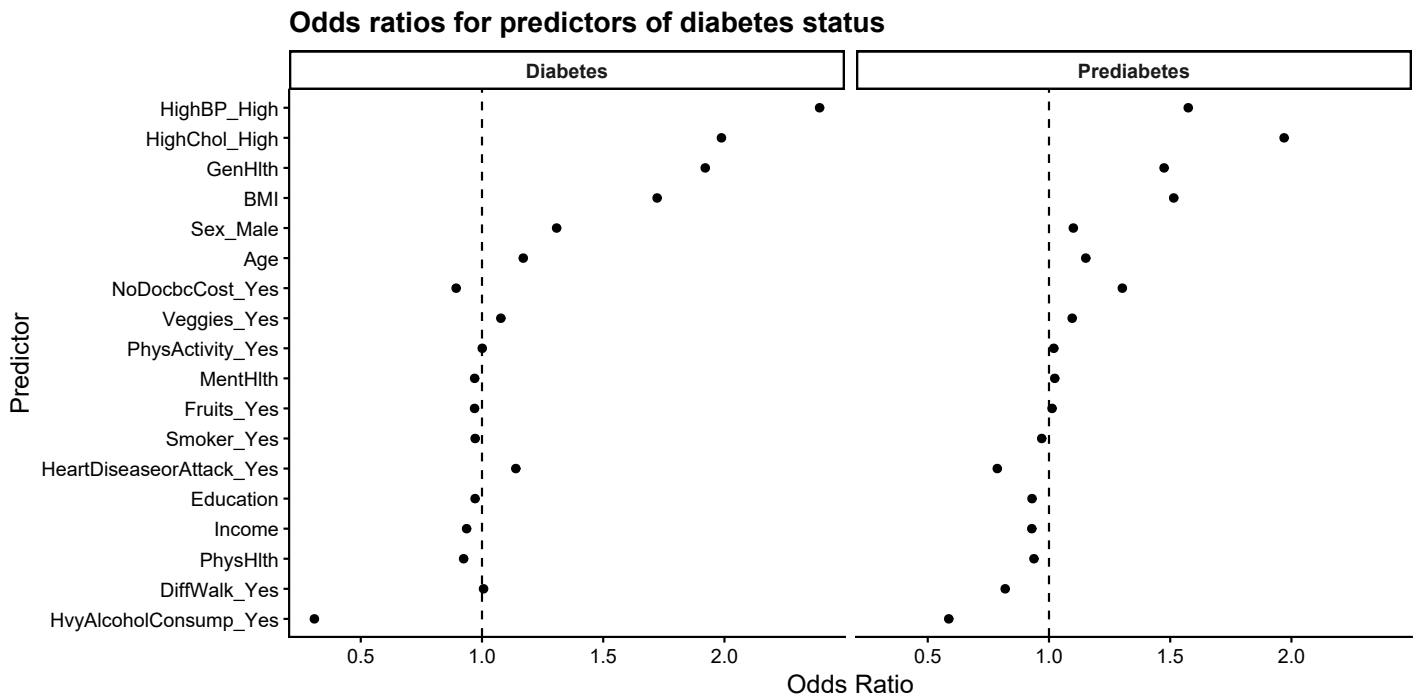

Multinomial Logistic Regression: Odds Ratio (OR) Summary

OR > 1 indicates increased odds; OR < 1 indicates decreased odds.

Synthesis of Findings

- **Primary Drivers:** Clinical indicators (Hypertension, Hypercholesterolemia, and BMI) represent the most significant risk. High blood pressure, in particular, yields a **139% increase** in the odds of diabetes ($OR = 2.39$).
- **The BMI Gradient:** The data shows a clear progression where BMI impacts diabetes ($OR = 1.72$) more severely than prediabetes ($OR = 1.51$), confirming its role in disease advancement.
- **Behavioral Mediation:** The lack of independent effect for physical activity ($OR = 1.00$) suggests its impact is likely **fully mediated** by metabolic variables like BMI and blood pressure.

- **Statistical Anomalies:** The negative association with heavy alcohol ($OR = 0.31$) likely reflects **selection bias** (e.g., "sick quitters" or health-conscious adjustments after diagnosis) rather than a biological protective effect.

Conclusion

Diabetes status in this population is best characterized by a **multifactorial cardiometabolic profile**. While age and sex are relevant, clinical markers of obesity and metabolic syndrome remain the most effective predictors for distinguishing between healthy, prediabetic, and diabetic individuals.

6.5.2 Global SHAP importance

```
sv_importance(sv, kind = "bar") +
  scale_fill_discrete(
    labels = c(
      "No diabetes",
      "Prediabetes",
      "Diabetes")
  )
```

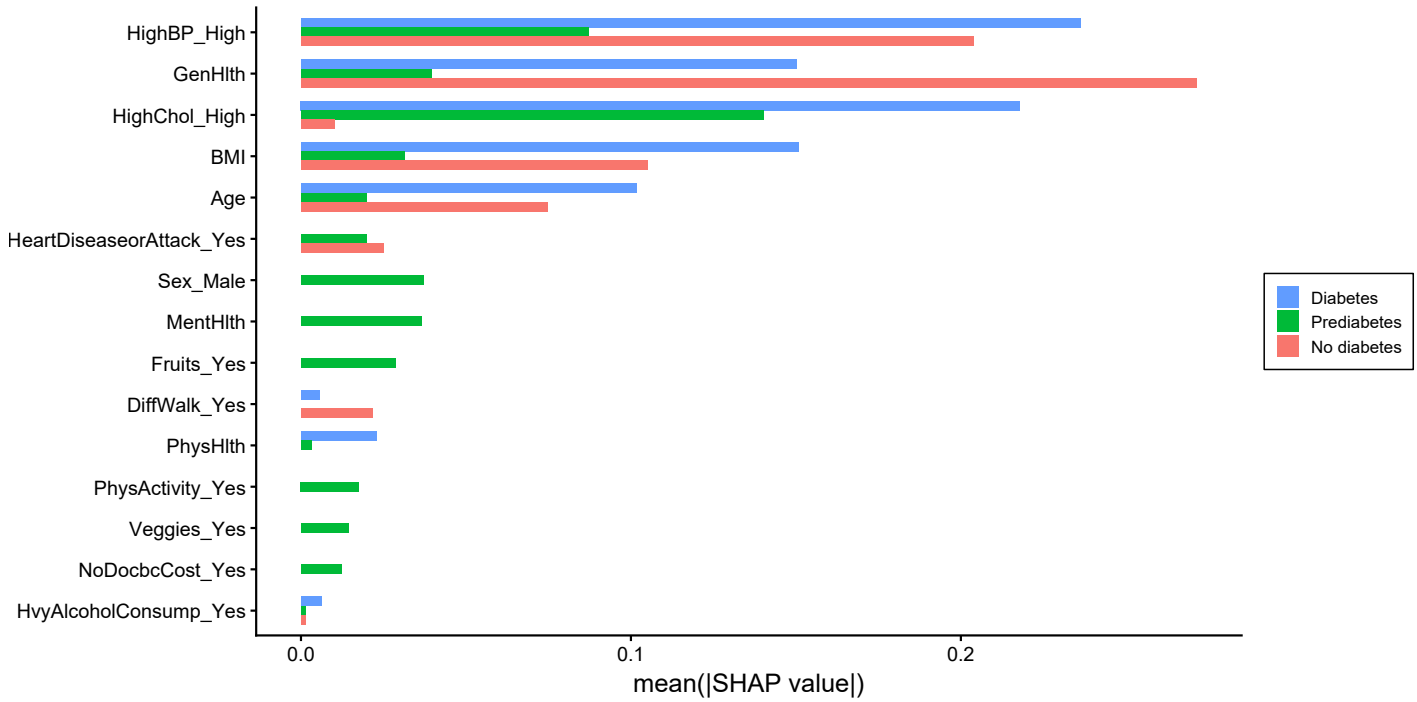

The SHAP importance plot indicates that cardiometabolic indicators dominate the predictive structure of the model. High blood pressure, general health status, high cholesterol, and BMI exhibit the largest mean absolute SHAP contributions across classes, indicating they play the strongest role in shifting predicted probabilities of diabetes status. Age also contributes meaningfully, while behavioural variables such as smoking, fruit consumption, and alcohol use show comparatively smaller effects. This pattern closely mirrors the multinomial logistic regression results, reinforcing the conclusion that diabetes status in this dataset is primarily driven by metabolic and health status indicators rather than lifestyle variables alone.

References

1. Aune, D., Norat, T., Leitzmann, M., Tonstad, S., & Vatten, L. J. (2015). Physical activity and the risk of type 2 diabetes: A systematic review and dose–response meta-analysis. *European Journal of Epidemiology*, 30(7), 529–542.
2. Hill-Briggs, F., Adler, N. E., Berkowitz, S. A., et al. (2021). Social determinants of health and diabetes: A scientific review. *Diabetes Care*, 44(1), 258–279.

3. Ismail, L., Materwala, H., & Al Kaabi, J. (2021). Association of risk factors with type 2 diabetes: A systematic review. *Computational and Structural Biotechnology Journal*, 19, 1759–1785.
4. Tatulashvili, S., Fagherazzi, G., Dow, C., Cohen, R., Fosse-Edorh, S., & Balkau, B. (2020). Socioeconomic inequalities and type 2 diabetes complications: A systematic review. *Diabetes & Metabolism*, 46(2), 89–99.
5. Liu, X., Faes, L., Kale, A. U., et al. (2022). Machine learning-based prediction models for cardiovascular disease risk using electronic health records data: systematic review and meta-analysis. *BMJ Open*, 12(6), e057995.
6. Lundberg, S. M., & Lee, S. I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 30, 4765–4774.
7. Roth, G. A., Johnson, C., Abajobir, A., et al. (2024). Heart disease and stroke statistics—2024 update: A report from the American Heart Association. *Circulation*, 149(8).
8. Collins, S. R., Gunja, M. Z., & Aboulafia, G. N. (2023). *Paying for it: How health care costs and medical debt are making Americans sicker and poorer*. Commonwealth Fund.
9. Kessler, R. C., Aguilar-Gaxiola, S., Alonso, J., et al. (2010). Relationship between household income and mental disorders. *Archives of General Psychiatry*, 67(5), 474–483.
10. Liu, J., Jiang, N., Fan, A. Z., Thompson, W. W., Ding, R., & Ni, S. (2023). Investigating the associations between socioeconomic factors and unhealthy days among adults using zero-inflated negative binomial regression. *SAGE Open*, 13(3).
11. Rust, K. F., Holt, J. B., & Lu, H. (2014). Models for count data with an application to healthy days measures. *Preventing Chronic Disease*, 11, E91.
12. Gunja, M. Z., et al. (2023). *The cost of not getting care: Income disparities in the affordability of health services across high-income countries*. Commonwealth Fund.